

# Openshift 4 Health Check Report

Red Hat Consulting



# Table of Contents

1. Commons	1
1.1. Versions	1
1.2. Components status	1
1.3. Cluster Status Conditions	1
1.4. Cluster Abnormal Events	1
1.5. Cluster API Server Status	1
1.6. Cluster Console Status	1
2. nodes	2
2.1. Cluster Failing Nodes	2
2.2. Cluster Node Conditions	2
2.3. Nodes Capacity	2
3. Machines	3
3.1. Machine Info	3
3.2. Machinesets Info	3
3.3. Machine Configs	3
3.4. Machine Configs	3
3.5. Machine Auto Scalers	3
3.6. Cluster Auto Scalers	3
3.7. Machine Health Checks	3
4. etcd	4
4.1. etcd pods info	4
4.2. Members in the cluster	4
4.3. Endpoints status	4
4.4. Endpoints Health	4
5. pods	5
5.1. Failing pods	5
5.2. Constantly restarted pods	5
5.3. Long running pods	5
5.4. Pod Disruption Budget	5
5.5. Pods in default namespace	6
5.6. Pods per node	6
6. performance	7
6.1. Nodes memory utilization	7
6.2. Nodes CPU utilization	7
6.3. Pods memory utilization	7
6.4. Pods CPU utilization	7

7. security	9
7.1. Pending CSRs	9
7.2. Cluster Identities	9
7.3. Identities grants	9
7.4. Rolebindings	10
7.5. Cluster Rolebindings	24
7.6. Kubeadmin Secret	29
7.7. Identity Providers	29
7.8. Cluster Authentications	29
7.9. Kubelet Configurations	30
7.10. Cluster Subscriptions	30
7.11. WebHook Configurations	30
7.12. Mutating WebHook Configurations	33
7.13. API Versions	34
8. storage	36
8.1. Failed Persistent Volumes	36
8.2. Pending Persistent Volumes	36
8.3. Released Persistent Volumes	36
8.4. Lost Persistent Volumes Claims	36
8.5. Pending Persistent Volumes Claims	36
8.6. Storage Classes	36
8.7. Quotas	36
8.8. Volume Snapshots	36
8.9. CSI Drivers	37
8.10. CSI Nodes	37
8.11. Feature Gates	37
8.12. Horizontal Pod AutoScalers	37
9. logging	38
9.1. Logging Resources	38
10. Monitoring	39
10.1. Prometheus Status	39
10.2. Prometheus Rules	39
10.3. Prometheus Context	77
10.4. Service Monitors	78
10.5. Pod Monitors	79
10.6. Alert Managers	79
10.7. Monitoring Agents & Dashboards	79
11. operators	80

11.1. Degraded Cluster Operators .....	80
11.2. Unavailable Cluster Operators .....	80
11.3. Cluster Services Versions .....	80
11.4. Operator Groups .....	80
11.5. Operator Sources .....	80
12. Service Mesh .....	81
12.1. Service Mesh ControlPlane .....	81
12.2. Service Mesh Members .....	81
12.3. Service Mesh Member Rolls .....	81
13. Applications .....	82
13.1. Non-Ready Deployments .....	82
13.2. Unavailable Deployments .....	82
13.3. Failed Builds .....	82
13.4. Inactive projects .....	82

# Chapter 1. Commons

## 1.1. Versions

```
Client Version: 4.10.18
Server Version: 4.10.18
Kubernetes Version: v1.23.5+3afdacb
```

## 1.2. Components status

NAME	STATUS	MESSAGE	ERROR
etcd-1	Healthy	{"health":"true","reason":""}	
etcd-0	Healthy	{"health":"true","reason":""}	
scheduler	Healthy	ok	
controller-manager	Healthy	ok	

## 1.3. Cluster Status Conditions

## 1.4. Cluster Abnormal Events

NAMESPACE	SEEN	TYPE	REASON
ocapp-test	Warning	FailedCreate	replicaset/my-app2-5f5c8df97d
ocapp-test	Warning	FailedCreate	replicaset/my-app3-7796c8c7bb

## 1.5. Cluster API Server Status

## 1.6. Cluster Console Status

## Chapter 2. nodes

### 2.1. Cluster Failing Nodes

### 2.2. Cluster Node Conditions

#### 2.2.1. master Nodes

**crc-5nvrn-master-0**

```
MemoryPressure False KubeletHasSufficientMemory  
DiskPressure False KubeletHasNoDiskPressure  
PIDPressure False KubeletHasSufficientPID
```

#### 2.2.2. worker Nodes

#### 2.2.3. infra Nodes

### 2.3. Nodes Capacity

NODE	TYPE	OS	CPU	MEM	STORAGE	PODS
crc-5nvrn-master-0	master	CoreOS	4	31970284Ki	15377152Ki	250

## Chapter 3. Machines

### 3.1. Machine Info

NAMESPACE	NAME	PHASE	TYPE	REGION	ZONE	AGE
openshift-machine-api	crc-5nvrn-master-0	Running				41d

### 3.2. Machinesets Info

NAMESPACE	NAME	DESIRED	CURRENT	READY	AVAILABLE	AGE
openshift-machine-api	crc-5nvrn-worker-0	0	0			41d

### 3.3. Machine Configs

### 3.4. Machine Configs

### 3.5. Machine Auto Scalers

### 3.6. Cluster Auto Scalers

### 3.7. Machine Health Checks

NAMESPACE	NAME	MAXUNHEALTHY	EXPECTEDMACHINES	CURRENTHEALTHY
openshift-machine-api	machine-api-termination-handler	100%	0	0

## Chapter 4. etcd

### 4.1. etcd pods info

NAME	READY	STATUS	RESTARTS	AGE
etcd-crc-5nvrn-master-0	4/4	Running	16	41d

### 4.2. Members in the cluster

ID	STATUS	NAME	PEER ADDRS	CLIENT ADDRS	IS LEARNER
6ef9bb7ffab9d73b	started	crc-5nvrn-master-0	<a href="https://192.168.126.11:2380">https://192.168.126.11:2380</a>	<a href="https://192.168.126.11:2379">https://192.168.126.11:2379</a>	false

### 4.3. Endpoints status

ENDPOINT	ID	VERSION	DB SIZE	IS LEADER	IS LEARNER	RAFT TERM	RAFT INDEX	RAFT APPLIED INDEX	ERRORS
<a href="https://192.168.126.11:2379">https://192.168.126.11:2379</a>	6ef9bb7ffab9d73b	3.5.3	86 MB	true	false	12	1881087	1881087	

### 4.4. Endpoints Health

ENDPOINT	HEALTH	TOOK	ERROR
<a href="https://192.168.126.11:2379">https://192.168.126.11:2379</a>	true	13.107493ms	



# Chapter 5. pods

## 5.1. Failing pods

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
ocapp-test	my-app2-1-build	0/1	Error	0	13d
ocapp-test	my-app3-1-build	0/1	Error	0	13d

## 5.2. Constanstly restarted pods

openshift-machine-api	machine-api-controllers-576c6d7795-vjzr5	7/7	Running	48	41d
openshift-kube-controller-manager	kube-controller-manager-crc-5nvrn-master-0	4/4	Running	24	41d
openshift-kube-apiserver	kube-apiserver-crc-5nvrn-master-0	5/5	Running	16	5d16h
openshift-etcd	etcd-crc-5nvrn-master-0	4/4	Running	16	41d
openshift-kube-scheduler	openshift-kube-scheduler-crc-5nvrn-master-0	3/3	Running	15	41d
openshift-cluster-machine-approver	machine-approver-697c6c67f-pl7qf	2/2	Running	15	(19h ago) 40d
openshift-sdn	sdn-controller-6fx72	2/2	Running	9	41d
openshift-operator-lifecycle-manager	packageserver-68868bbf8c-jpxgm	1/1	Running	9	41d
openshift-machine-api	machine-api-operator-5869dbb6f9-65zc2	2/2	Running	9	41d
openshift-sdn	sdn-552c8	2/2	Running	8	41d

## 5.3. Long running pods

openshift-marketplace	redhat-operators-6t7vs	1/1	Running	0	143m
openshift-service-ca	service-ca-7567987bc8-7wl9j	1/1	Running	5	41d
openshift-sdn	sdn-controller-6fx72	2/2	Running	9	41d
openshift-sdn	sdn-552c8	2/2	Running	8	41d
openshift-operator-lifecycle-manager	packageserver-68868bbf8c-jpxgm	1/1	Running	9	41d
openshift-oauth-apiserver	apiserver-65b8444f69-zfxrz	1/1	Running	4	41d
openshift-network-diagnostics	network-check-target-85xvd	1/1	Running	4	41d
openshift-network-diagnostics	network-check-source-695cdb7747-cf742	1/1	Running	4	41d
openshift-multus	network-metrics-daemon-5gfkl	2/2	Running	8	41d
openshift-multus	multus-bzhmc	1/1	Running	4	41d

## 5.4. Pod Disruption Budget

NAMESPACE	NAME	MIN	AVAILABLE	MAX	UNAVAILABLE	ALLOWED
DISRUPTIONS	AGE					
openshift-image-registry	image-registry	0	N/A	1	41d	
openshift-operator-lifecycle-manager	packageserver-pdb	N/A	1	1	41d	

## 5.5. Pods in default namespace

## 5.6. Pods per node

88 crc-5nvrn-master-0

# Chapter 6. performance

## 6.1. Nodes memory utilization

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
crc-5nvrn-master-0	2158m	56%	10591Mi	72%

## 6.2. Nodes CPU utilization

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
crc-5nvrn-master-0	2158m	56%	10591Mi	72%

## 6.3. Pods memory utilization

NAMESPACE	NAME	CPU(cores)	
MEMORY(bytes)			
openshift-kube-apiserver	kube-apiserver-crc-5nvrn-master-0	272m	1748Mi
openshift-etcd	etcd-crc-5nvrn-master-0	190m	1209Mi
openshift-monitoring	prometheus-k8s-0	121m	1095Mi
openshift-kube-controller-manager	kube-controller-manager-crc-5nvrn-master-0	20m	447Mi
openshift-apiserver	apiserver-6c7d78dfc4-v28gn	47m	246Mi
openshift-machine-api	machine-api-controllers-576c6d7795-vjzr5	1m	176Mi
openshift-operator-lifecycle-manager	packageserver-68868bbf8c-jpxgm	70m	156Mi
openshift-authentication-operator	authentication-operator-6ddbdbf46-9rsc7	58m	153Mi
openshift-marketplace	certified-operators-fztjt	7m	148Mi
openshift-controller-manager	controller-manager-tdftp	5m	147Mi
openshift-monitoring	thanos-querier-647cd76596-tx86p	43m	141Mi
openshift-operator-lifecycle-manager	olm-operator-67fcd489c-96kxh	3m	138Mi
openshift-marketplace	community-operators-6lvjd	15m	136Mi
openshift-service-ca	service-ca-7567987bc8-7w19j	4m	132Mi
openshift-ingress-operator	ingress-operator-7899578f6-b5ljs	2m	132Mi
openshift-monitoring	prometheus-operator-76f79c8d85-xn57s	5m	130Mi
openshift-kube-scheduler	openshift-kube-scheduler-crc-5nvrn-master-0	4m	125Mi
openshift-kube-apiserver-operator	kube-apiserver-operator-5f76847f99-8svn8	11m	123Mi
openshift-network-operator	network-operator-7d77dffff9-qfzzx	5m	120Mi
openshift-monitoring	grafana-5487f6d9cf-vxp6k	8m	112Mi

## 6.4. Pods CPU utilization

NAMESPACE	NAME	CPU(cores)	
MEMORY(bytes)			
openshift-kube-apiserver	kube-apiserver-crc-5nvrn-master-0	272m	1748Mi
openshift-etcd	etcd-crc-5nvrn-master-0	190m	1209Mi
openshift-monitoring	prometheus-k8s-0	121m	1095Mi
openshift-operator-lifecycle-manager	packageserver-68868bbf8c-jpxgm	70m	156Mi
openshift-authentication-operator	authentication-operator-6ddbdbbf46-9rsc7	58m	153Mi
openshift-apiserver	apiserver-6c7d78dfc4-v28gn	47m	246Mi
openshift-monitoring	thanos-querier-647cd76596-tx86p	43m	141Mi
openshift-oauth-apiserver	apiserver-65b8444f69-zfxrz	37m	76Mi
openshift-kube-controller-manager	kube-controller-manager-crc-5nvrn-master-0	20m	447Mi
openshift-multus	multus-bzhmc	16m	6Mi
openshift-marketplace	community-operators-6lvjd	15m	136Mi
openshift-monitoring	prometheus-adapter-6dbdddc76b-bdr2v	14m	56Mi
openshift-kube-apiserver-operator	kube-apiserver-operator-5f76847f99-8svn8	11m	123Mi
openshift-etcd-operator	etcd-operator-6d6ff5499d-cd588	10m	74Mi
openshift-marketplace	redhat-operators-6t7vs	10m	67Mi
openshift-cluster-samples-operator	cluster-samples-operator-b6fc87b94-w65rt	9m	80Mi
openshift-apiserver-operator	openshift-apiserver-operator-6756c64b88-kh9qj	9m	110Mi
openshift-monitoring	grafana-5487f6d9cf-vxp6k	8m	112Mi
openshift-console-operator	console-operator-c74b55d67-9mskj	8m	71Mi
openshift-marketplace	redhat-marketplace-bhgkk	7m	90Mi

# Chapter 7. security

## 7.1. Pending CSRs

## 7.2. Cluster Identities

NAME	IDP NAME	IDP USER NAME	USER NAME	USER UID
developer:developer	developer	developer	developer	ca077f0b-2f5e-4304-9c32-8046d4ede025
developer:kubeadmin	developer	kubeadmin	kubeadmin	8facc26a-4157-4d65-be96-ee5b50ff2b0d

## 7.3. Identities grants

### 7.3.1. Identities who can create users

```
Users: kubeadmin
       system:admin
       system:serviceaccount:openshift-apiserver-operator:openshift-apiserver-operator
       system:serviceaccount:openshift-apiserver:openshift-apiserver-sa
       system:serviceaccount:openshift-authentication-operator:authentication-operator
       system:serviceaccount:openshift-authentication:oauth-openshift
       system:serviceaccount:openshift-cluster-storage-operator:cluster-storage-operator
       system:serviceaccount:openshift-cluster-storage-operator:csi-snapshot-controller-operator
       system:serviceaccount:openshift-cluster-version:default
       system:serviceaccount:openshift-config-operator:openshift-config-operator
       system:serviceaccount:openshift-controller-manager-operator:openshift-controller-manager-operator
       system:serviceaccount:openshift-etcd-operator:etcd-operator
       system:serviceaccount:openshift-etcd:installer-sa
       system:serviceaccount:openshift-kube-apiserver-operator:kube-apiserver-operator
       system:serviceaccount:openshift-kube-apiserver:installer-sa
       system:serviceaccount:openshift-kube-apiserver:localhost-recovery-client
       system:serviceaccount:openshift-kube-controller-manager-operator:kube-controller-manager-operator
       system:serviceaccount:openshift-kube-controller-manager:installer-sa
       system:serviceaccount:openshift-kube-controller-manager:localhost-recovery-client
       system:serviceaccount:openshift-kube-scheduler-operator:openshift-kube-scheduler-operator
       system:serviceaccount:openshift-kube-scheduler:installer-sa
       system:serviceaccount:openshift-kube-scheduler:localhost-recovery-client
       system:serviceaccount:openshift-kube-storage-version-migrator-operator:kube-storage-version-
migrator-operator
       system:serviceaccount:openshift-machine-config-operator:default
       system:serviceaccount:openshift-network-operator:default
       system:serviceaccount:openshift-oauth-apiserver:oauth-apiserver-sa
       system:serviceaccount:openshift-operator-lifecycle-manager:olm-operator-serviceaccount
       system:serviceaccount:openshift-service-ca-operator:service-ca-operator
```

## 7.3.2. Identities who can delete users

```

Users: kubeadmin
      system:admin
      system:serviceaccount:kube-system:generic-garbage-collector
      system:serviceaccount:kube-system:namespace-controller
      system:serviceaccount:openshift-apiserver-operator:openshift-apiserver-operator
      system:serviceaccount:openshift-apiserver:openshift-apiserver-sa
      system:serviceaccount:openshift-authentication-operator:authentication-operator
      system:serviceaccount:openshift-authentication:oauth-openshift
      system:serviceaccount:openshift-cluster-storage-operator:cluster-storage-operator
      system:serviceaccount:openshift-cluster-storage-operator:csi-snapshot-controller-operator
      system:serviceaccount:openshift-cluster-version:default
      system:serviceaccount:openshift-config-operator:openshift-config-operator
      system:serviceaccount:openshift-controller-manager-operator:openshift-controller-manager-operator
      system:serviceaccount:openshift-etcd-operator:etcd-operator
      system:serviceaccount:openshift-etcd:installer-sa
      system:serviceaccount:openshift-kube-apiserver-operator:kube-apiserver-operator
      system:serviceaccount:openshift-kube-apiserver:installer-sa
      system:serviceaccount:openshift-kube-apiserver:localhost-recovery-client
      system:serviceaccount:openshift-kube-controller-manager-operator:kube-controller-manager-operator
      system:serviceaccount:openshift-kube-controller-manager:installer-sa
      system:serviceaccount:openshift-kube-controller-manager:localhost-recovery-client
      system:serviceaccount:openshift-kube-scheduler-operator:openshift-kube-scheduler-operator
      system:serviceaccount:openshift-kube-scheduler:installer-sa
      system:serviceaccount:openshift-kube-scheduler:localhost-recovery-client
      system:serviceaccount:openshift-kube-storage-version-migrator-operator:kube-storage-version-
migrator-operator
      system:serviceaccount:openshift-machine-config-operator:default
      system:serviceaccount:openshift-network-operator:default
      system:serviceaccount:openshift-oauth-apiserver:oauth-apiserver-sa
      system:serviceaccount:openshift-operator-lifecycle-manager:olm-operator-serviceaccount
      system:serviceaccount:openshift-service-ca-operator:service-ca-operator

```

## 7.4. Rolebindings

default

NAME	ROLE
machine-config-controller-events	ClusterRole/machine-config-controller-events
machine-config-daemon-events	ClusterRole/machine-config-daemon-events
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

delete

NAME	ROLE
admin	ClusterRole/admin
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### kube-node-lease

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### kube-public

NAME	ROLE
system:controller:bootstrap-signer	Role/system:controller:bootstrap-signer
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### kube-system

NAME	ROLE
authentication-reader-for-authenticated-users-reader	Role/extension-apiserver-authentication-
cloud-controller-manager:apiserver-authentication-reader	Role/extension-apiserver-authentication-
console-reader	Role/extension-apiserver-authentication-
console-operator-reader	Role/extension-apiserver-authentication-
insights-operator-auth-reader	Role/extension-apiserver-authentication-
network-diagnostics-reader	Role/extension-apiserver-authentication-
packageserver-service-auth-reader	Role/extension-apiserver-authentication-
prometheus-k8s	Role/prometheus-k8s
resource-metrics-auth-reader	Role/extension-apiserver-authentication-
system::extension-apiserver-authentication-reader	Role/extension-apiserver-authentication-
system::leader-locking-kube-controller-manager-controller-manager	Role/system::leader-locking-kube-
system::leader-locking-kube-scheduler	Role/system::leader-locking-kube-scheduler
system:controller:bootstrap-signer	Role/system:controller:bootstrap-signer
system:controller:cloud-provider	Role/system:controller:cloud-provider
system:controller:token-cleaner	Role/system:controller:token-cleaner
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:controller:kube-apiserver-check-endpoints-reader	Role/extension-apiserver-authentication-
system:openshift:leader-election-lock-kube-controller-manager-kube-controller-manager	Role/system:openshift:leader-election-lock-
system:openshift:leader-locking-kube-controller-manager-controller-manager	Role/system::leader-locking-kube-
system:openshift:leader-locking-kube-scheduler	Role/system::leader-locking-kube-scheduler
system:openshift:leader-locking-openshift-controller-manager-openshift-controller-manager	Role/system:openshift:leader-locking-

## ocapp-test

NAME	ROLE
admin	ClusterRole/admin
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift



NAME	ROLE
cluster-samples-operator-openshift-edit	ClusterRole/edit
shared-resource-viewers	Role/shared-resource-viewer
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-apiserver

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-apiserver-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-authentication

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-authentication-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cloud-controller-manager

NAME	ROLE
cloud-controller-manager	Role/cloud-controller-manager
cluster-cloud-controller-manager	ClusterRole/admin
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cloud-controller-manager-operator

NAME	ROLE
cluster-cloud-controller-manager	Role/cluster-cloud-controller-manager
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cloud-credential-operator

NAME	ROLE
cloud-credential-operator	Role/cloud-credential-operator-role
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cloud-network-config-controller

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cluster-csi-drivers

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cluster-machine-approver

NAME	ROLE
machine-approver	Role/machine-approver
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cluster-node-tuning-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cluster-samples-operator

NAME	ROLE
cluster-samples-operator	Role/cluster-samples-operator
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cluster-storage-operator

NAME	ROLE
csi-snapshot-controller-leader-election	Role/csi-snapshot-controller-leader-election
prometheus	Role/prometheus
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-cluster-version

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-config

NAME	ROLE
cluster-cloud-controller-manager	Role/cluster-cloud-controller-manager
cluster-samples-operator-openshift-config-secret-reader	Role/coreos-pull-secret-reader
console-operator	Role/console-operator
ingress-operator	Role/ingress-operator
insights-operator	Role/insights-operator
machine-api-controllers	Role/machine-api-controllers
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-config-managed

NAME	ROLE
cluster-cloud-controller-manager	Role/cluster-cloud-controller-manager
console-configmap-reader	Role/console-configmap-reader
console-operator	Role/console-operator
console-public	Role/console-public
insights-operator-etc-pki-entitlement	Role/insights-operator-etc-pki-entitlement
machine-api-controllers	Role/machine-api-controllers
machine-approver	Role/machine-approver
openshift-network-public-role-binding	Role/openshift-network-public-role
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:oauth-servercert-trust	Role/system:openshift:oauth-servercert-trust

#### openshift-config-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-console

NAME	ROLE
console-operator	Role/console-operator
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-console-operator

NAME	ROLE
console-operator	Role/console-operator
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-console-user-settings

NAME	ROLE
console-user-settings-admin	Role/console-user-settings-admin
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
user-settings-8facc26a-4157-4d65-be96-ee5b50ff2b0d-rolebinding	Role/user-settings-8facc26a-4157-4d65-be96-ee5b50ff2b0d-role
user-settings-ca077f0b-2f5e-4304-9c32-8046d4ede025-rolebinding	Role/user-settings-ca077f0b-2f5e-4304-9c32-8046d4ede025-role

## openshift-controller-manager

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:leader-locking-openshift-controller-manager	Role/system:openshift:leader-locking-openshift-controller-manager

## openshift-controller-manager-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-dns

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-dns-operator

NAME	ROLE
dns-operator	Role/dns-operator
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-etcd

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-etcd-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-host-network

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-image-registry

NAME	ROLE
cluster-image-registry-operator	Role/cluster-image-registry-operator
node-ca	Role/node-ca
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-infra

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:sa-creating-openshift-controller-manager	Role/system:openshift:sa-creating-openshift-controller-manager

#### openshift-ingress

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-ingress-canary

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-ingress-operator

NAME	ROLE
ingress-operator	Role/ingress-operator
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-insights

NAME	ROLE
insights-operator	Role/insights-operator
insights-operator-obfuscation-secret	Role/insights-operator-obfuscation-secret
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-kni-infra

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-kube-apiserver

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:controller:check-endpoints	ClusterRole/system:openshift:controller:check-endpoints

## openshift-kube-apiserver-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-kube-controller-manager

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:leader-election-lock-cluster-policy-controller	Role/system:openshift:leader-election-lock-cluster-policy-controller

#### openshift-kube-controller-manager-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-kube-scheduler

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:sa-listing-configmaps	Role/system:openshift:sa-listing-configmaps

#### openshift-kube-scheduler-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-kube-storage-version-migrator-operator

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller



## openshift-machine-api

NAME	ROLE
cluster-autoscaler	Role/cluster-autoscaler
cluster-autoscaler-operator	Role/cluster-autoscaler-operator
cluster-baremetal-operator	Role/cluster-baremetal-operator
machine-api-controllers	Role/machine-api-controllers
machine-api-operator	Role/machine-api-operator
prometheus-k8s-cluster-autoscaler-operator	Role/prometheus-k8s-cluster-autoscaler-operator
prometheus-k8s-cluster-baremetal-operator	Role/prometheus-k8s-cluster-baremetal-operator
prometheus-k8s-machine-api-operator	Role/prometheus-k8s-machine-api-operator
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-machine-config-operator

NAME	ROLE
machine-config-controller-events	ClusterRole/machine-config-controller-events
machine-config-daemon-events	ClusterRole/machine-config-daemon-events
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-marketplace

NAME	ROLE
marketplace-operator	Role/marketplace-operator
openshift-marketplace-metrics	Role/openshift-marketplace-metrics
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-monitoring

NAME	ROLE
alertmanager-prometheusk8s	Role/monitoring-alertmanager-edit
cluster-monitoring-operator	ClusterRole/cluster-monitoring-operator-namespaced
insights-operator-alertmanager	Role/monitoring-alertmanager-edit
prometheus-k8s	Role/prometheus-k8s
prometheus-k8s-config	Role/prometheus-k8s-config
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-multus

NAME	ROLE
multus-whereabouts	Role/whereabouts-cni
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-network-diagnostics

NAME	ROLE
network-diagnostics	Role/network-diagnostics
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-network-operator

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-node

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:node-config-reader	Role/system:node-config-reader

#### openshift-oauth-apiserver

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-openstack-infra

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-operator-lifecycle-manager

NAME	ROLE
collect-profiles	Role/collect-profiles
operator-lifecycle-manager-metrics	Role/operator-lifecycle-manager-metrics
packageserver	Role/packageserver
packageserver-service-cert	Role/packageserver-service-cert
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-operators

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-ovirt-infra

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-sdn

NAME	ROLE
openshift-sdn-controller-leaderelection	Role/openshift-sdn-controller-leaderelection
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## openshift-service-ca

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller
system:openshift:controller:service-ca	Role/system:openshift:controller:service-ca

## openshift-service-ca-operator

NAME	ROLE
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-user-workload-monitoring

NAME	ROLE
cluster-monitoring-operator	ClusterRole/cluster-monitoring-operator-namespaced
prometheus-k8s	Role/prometheus-k8s
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

#### openshift-vsphere-infra

NAME	ROLE
system:deployers	ClusterRole/system:deployer
system:image-builders	ClusterRole/system:image-builder
system:image-pullers	ClusterRole/system:image-puller

## 7.5. Cluster Rolebindings

system:openshift:controller:template-instance-controller:admin	admin
system:openshift:controller:template-instance-finalizer-controller:admin	admin
alertmanager-main	alertmanager-main
basic-users	basic-user
cloud-controller-manager	cloud-controller-manager
cloud-credential-operator-rolebinding	cloud-credential-operator-role
cloud-node-manager	cloud-node-manager
cluster-admin	cluster-admin
cluster-admins	cluster-admin
cluster-storage-operator-role	cluster-admin
cluster-version-operator	cluster-admin
csi-snapshot-controller-operator-role	cluster-admin
default-account-cluster-network-operator	cluster-admin
default-account-openshift-machine-config-operator	cluster-admin
kubeadmin	cluster-admin
system:openshift:oauth-apiserver	cluster-admin
system:openshift:openshift-apiserver	cluster-admin
system:openshift:openshift-authentication	cluster-admin
system:openshift:operator:authentication	cluster-admin
system:openshift:operator:cluster-kube-scheduler-operator	cluster-admin
system:openshift:operator:etcd-operator	cluster-admin
system:openshift:operator:kube-apiserver-operator	cluster-admin
system:openshift:operator:kube-apiserver-recovery	cluster-admin
system:openshift:operator:kube-controller-manager-operator	cluster-admin
system:openshift:operator:kube-controller-manager-recovery	cluster-admin
system:openshift:operator:kube-scheduler-recovery	cluster-admin

system:openshift:operator:kube-storage-version-migrator-operator	cluster-admin
system:openshift:operator:openshift-apiserver-operator	cluster-admin
system:openshift:operator:openshift-config-operator	cluster-admin
system:openshift:operator:openshift-controller-manager-operator	cluster-admin
system:openshift:operator:openshift-etcd-installer	cluster-admin
system:openshift:operator:openshift-kube-apiserver-installer	cluster-admin
system:openshift:operator:openshift-kube-controller-manager-installer	cluster-admin
system:openshift:operator:openshift-kube-scheduler-installer	cluster-admin
system:openshift:operator:service-ca-operator	cluster-admin
cluster-autoscaler	cluster-admin
cluster-autoscaler-operator	cluster-autoscaler
cluster-baremetal-operator	cluster-autoscaler-operator
default-account-cluster-image-registry-operator	cluster-baremetal-operator
cluster-monitoring-operator	cluster-image-registry-operator
prometheus-adapter-view	cluster-monitoring-operator
telemeter-client-view	cluster-monitoring-view
cluster-node-tuning-operator	cluster-monitoring-view
cluster-node-tuning:tuned	cluster-node-tuning-operator
cluster-readers	cluster-node-tuning:tuned
insights-operator-gather-reader	cluster-reader
cluster-samples-operator	cluster-reader
cluster-samples-operator-imageconfig-reader	cluster-samples-operator
imageconfig-reader	cluster-samples-operator-
cluster-samples-operator-proxy-reader	cluster-samples-operator-proxy-
reader	
cluster-status-binding	cluster-status
console	console
console-extensions-reader	console-extensions-reader
console-operator	console-operator
dns-monitoring	dns-monitoring
grafana	grafana
helm-chartrepos-view	helm-chartrepos-viewer
insights-operator	insights-operator
insights-operator-gather	insights-operator-gather
kube-apiserver	kube-apiserver
kube-state-metrics	kube-state-metrics
machine-api-controllers	machine-api-controllers
machine-api-operator	machine-api-operator
machine-api-operator-ext-remediation	machine-api-operator-ext-
remediation	
machine-config-controller	machine-config-controller
machine-config-daemon	machine-config-daemon
machine-config-server	machine-config-server
marketplace-operator	marketplace-operator
metrics-daemon-sa-rolebinding	metrics-daemon-role
multus	multus
multus-admission-controller-webhook	multus-admission-controller-
webhook	
network-diagnostics	network-diagnostics
node-exporter	node-exporter
openshift-csi-snapshot-controller-role	openshift-csi-snapshot-
controller-runner	
openshift-dns	openshift-dns
openshift-dns-operator	openshift-dns-operator
openshift-ingress-operator	openshift-ingress-operator
openshift-ingress-router	openshift-ingress-router
openshift-sdn	openshift-sdn

openshift-sdn-controller	openshift-sdn-controller
openshift-state-metrics	openshift-state-metrics
prometheus-adapter	prometheus-adapter
prometheus-k8s	prometheus-k8s
prometheus-k8s-scheduler-resources	prometheus-k8s-scheduler-
resources	
prometheus-operator	prometheus-operator
registry-monitoring	registry-monitoring
NAME	ROLE
router-monitoring	router-monitoring
self-access-reviewers	self-access-reviewer
self-provisioners	self-provisioner
console-auth-delegator	system:auth-delegator
console-operator-auth-delegator	system:auth-delegator
insights-operator-auth	system:auth-delegator
packageserver-service-system:auth-delegator	system:auth-delegator
resource-metrics:system:auth-delegator	system:auth-delegator
system:openshift:controller:kube-apiserver-check-endpoints-auth-delegator	system:auth-delegator
system:basic-user	system:basic-user
system:build-strategy-docker-binding	system:build-strategy-docker
system:build-strategy-jenkinspipeline-binding	system:build-strategy-
jenkinspipeline	
system:build-strategy-source-binding	system:build-strategy-source
system-bootstrap-node-renewal	
system:certificates.k8s.io:certificatesigningrequests:selfnodeclient	
system:controller:attachdetach-controller	system:controller:attachdetach-
controller	
system:controller:certificate-controller	system:controller:certificate-
controller	
system:controller:clusterrole-aggregation-controller	system:controller:clusterrole-
aggregation-controller	
system:controller:cronjob-controller	system:controller:cronjob-
controller	
system:controller:daemon-set-controller	system:controller:daemon-set-
controller	
system:controller:deployment-controller	system:controller:deployment-
controller	
system:controller:disruption-controller	system:controller:disruption-
controller	
system:controller:endpoint-controller	system:controller:endpoint-
controller	
system:controller:endpointslice-controller	
system:controller:endpointslice-controller	
system:controller:endpointslicemirroring-controller	
system:controller:endpointslicemirroring-controller	
system:controller:ephemeral-volume-controller	system:controller:ephemeral-
volume-controller	
system:controller:expand-controller	system:controller:expand-
controller	
system:controller:generic-garbage-collector	system:controller:generic-
garbage-collector	
system:controller:horizontal-pod-autoscaler	system:controller:horizontal-
pod-autoscaler	
system:controller:job-controller	system:controller:job-
controller	
system:controller:namespace-controller	system:controller:namespace-
controller	

system:controller:node-controller	system:controller:node-
controller	
olm-operator-binding-openshift-operator-lifecycle-manager	system:controller:operator-
lifecycle-manager	
system:controller:persistent-volume-binder	system:controller:persistent-
volume-binder	
system:controller:pod-garbage-collector	system:controller:pod-garbage-
collector	
system:controller:pvc-protection-controller	system:controller:pvc-
protection-controller	
system:controller:pv-protection-controller	system:controller:pv-
protection-controller	
system:controller:replicaset-controller	system:controller:replicaset-
controller	
system:controller:replication-controller	system:controller:replication-
controller	
system:controller:resourcequota-controller	
system:controller:resourcequota-controller	
system:controller:root-ca-cert-publisher	system:controller:root-ca-cert-
publisher	
system:controller:route-controller	system:controller:route-
controller	
system:controller:service-account-controller	system:controller:service-
account-controller	
system:controller:service-ca-cert-publisher	system:controller:service-ca-
cert-publisher	
system:controller:service-controller	system:controller:service-
controller	
system:controller:statefulset-controller	system:controller:statefulset-
controller	
system:controller:ttl-after-finished-controller	system:controller:ttl-after-
finished-controller	
system:controller:ttl-controller	system:controller:ttl-
controller	
system:deployer	system:deployer
system:discovery	system:discovery
system:image-builder	system:image-builder
openshift-image-registry-pruner	system:image-pruner
system:image-puller	system:image-puller
system:kube-controller-manager	system:kube-controller-manager
system:kube-dns	system:kube-dns
system:kube-scheduler	system:kube-scheduler
system:openshift:operator:kube-scheduler:public-2	system:kube-scheduler
system:masters	system:master
system:monitoring	system:monitoring
system:node	system:node
system:node-admins	system:node-admin
system:node-admin	system:node-admin
system-bootstrap-node-bootstrapper	system:node-bootstrapper
system:node-bootstrapper	system:node-bootstrapper
system:node-proxiers	system:node-proxier
system:node-proxier	system:node-proxier
system:oauth-token-deleters	system:oauth-token-deleter
system:openshift:controller:build-config-change-controller	
system:openshift:controller:build-config-change-controller	
system:openshift:controller:build-controller	
system:openshift:controller:build-controller	

```

system:openshift:controller:kube-apiserver-check-endpoints-crd-reader
system:openshift:controller:check-endpoints-crd-reader
system:openshift:controller:kube-apiserver-check-endpoints-node-reader
system:openshift:controller:check-endpoints-node-reader
system:openshift:controller:cluster-csr-approver-controller
system:openshift:controller:cluster-csr-approver-controller
system:openshift:controller:cluster-quota-reconciliation-controller
system:openshift:controller:cluster-quota-reconciliation-controller
system:openshift:controller:default-rolebindings-controller
system:openshift:controller:default-rolebindings-controller
system:openshift:controller:deployer-controller
system:openshift:controller:deployer-controller
system:openshift:controller:deploymentconfig-controller
system:openshift:controller:deploymentconfig-controller
system:openshift:controller:horizontal-pod-autoscaler
system:openshift:controller:horizontal-pod-autoscaler
system:openshift:controller:image-import-controller
system:openshift:controller:image-import-controller
system:openshift:controller:image-trigger-controller
system:openshift:controller:image-trigger-controller
system:openshift:controller:machine-approver
system:openshift:controller:machine-approver
system:openshift:controller:namespace-security-allocation-controller
system:openshift:controller:namespace-security-allocation-controller
system:openshift:controller:origin-namespace-controller
system:openshift:controller:origin-namespace-controller
system:openshift:controller:pv-recycler-controller
system:openshift:controller:pv-recycler-controller
system:openshift:controller:resourcequota-controller
system:openshift:controller:resourcequota-controller
system:openshift:controller:serviceaccount-controller
system:openshift:controller:serviceaccount-controller
system:openshift:controller:serviceaccount-pull-secrets-controller
system:openshift:controller:serviceaccount-pull-secrets-controller
system:openshift:controller:service-ca
system:openshift:controller:service-ca
system:openshift:controller:service-ingress-ip-controller
system:openshift:controller:service-ingress-ip-controller
system:openshift:controller:service-serving-cert-controller
system:openshift:controller:service-serving-cert-controller
system:openshift:controller:template-instance-controller
system:openshift:controller:template-instance-controller
system:openshift:controller:template-instance-finalizer-controller
system:openshift:controller:template-instance-finalizer-controller
system:openshift:controller:template-service-broker
system:openshift:controller:template-service-broker
system:openshift:controller:unidling-controller
system:openshift:controller:unidling-controller
system:openshift:discovery
system:openshift:kube-controller-manager:gce-cloud-provider
system:openshift:kube-controller-manager:gce-cloud-provider
system:openshift:openshift-controller-manager
system:openshift:openshift-controller-manager
system:openshift:openshift-controller-manager:ingress-to-route-controller
system:openshift:openshift-controller-manager:ingress-to-route-controller
system:openshift:openshift-controller-manager:update-buildconfig-status
system:openshift:openshift-controller-manager:update-buildconfig-status
system:openshift:controller:pv-
system:openshift:discovery
system:openshift:kube-
system:openshift:openshift-
system:openshift:openshift-
system:openshift:openshift-

```



```
system:openshift:operator:cloud-controller-manager
system:openshift:operator:cloud-controller-manager
system:openshift:public-info-viewer
system:openshift:tokenreview-openshift-controller-manager
system:openshift:tokenreview-openshift-controller-manager
system:openshift:useroauthaccesstoken-manager
system:openshift:useroauthaccesstoken-manager
system:public-info-viewer
registry-registry-role
system:scope-impersonation
system:sdn-readers
system:service-account-issuer-discovery
system:service-account-issuer-discovery
system:volume-scheduler
system:webhooks
telemeter-client
thanos-querier
multus-whereabouts
system:openshift:public-info-viewer
system:openshift:tokenreview-openshift-controller-manager
system:public-info-viewer
system:registry
system:scope-impersonation
system:sdn-reader
system:service-account-issuer-discovery
system:volume-scheduler
system:webhook
telemeter-client
thanos-querier
whereabouts-cni
```

## 7.6. Kubeadmin Secret

## 7.7. Identity Providers

```
[
  {
    "htpasswd": {
      "fileData": {
        "name": "htpass-secret"
      }
    },
    "mappingMethod": "claim",
    "name": "developer",
    "type": "HTPasswd"
  }
]
```

## 7.8. Cluster Authentications

```
{
  "include.release.openshift.io/ibm-cloud-managed": "true",
  "include.release.openshift.io/self-managed-high-availability": "true",
  "include.release.openshift.io/single-node-developer": "true",
  "release.openshift.io/create-only": "true"
}
{
  "kubeConfig": {
    "name": "webhook-authentication-integrated-oauth"
  }
}
{
  "integratedOAuthMetadata": {
    "name": "oauth-openshift"
  }
}
}
```

## 7.9. Kubelet Configurations

```
{
  "apiVersion": "v1",
  "items": [],
  "kind": "List",
  "metadata": {
    "resourceVersion": "",
    "selfLink": ""
  }
}
```

## 7.10. Cluster Subscriptions

## 7.11. WebHook Configurations

```
{
  "admissionReviewVersions": [
    "v1"
  ],
  "clientConfig": {
    "service": {
      "name": "machine-api-operator-webhook",
      "namespace": "openshift-machine-api",
      "path": "/validate-machine-openshift-io-v1beta1-machine",
      "port": 443
    }
  },
  "failurePolicy": "Ignore",
  "matchPolicy": "Equivalent",
}
```

```
"name": "validation.machine.machine.openshift.io",
"namespaceSelector": {},
"objectSelector": {},
"rules": [
  {
    "apiGroups": [
      "machine.openshift.io"
    ],
    "apiVersions": [
      "v1beta1"
    ],
    "operations": [
      "CREATE",
      "UPDATE"
    ],
    "resources": [
      "machines"
    ],
    "scope": "*"
  }
],
"sideEffects": "None",
"timeoutSeconds": 10
}
{
  "admissionReviewVersions": [
    "v1"
  ],
  "clientConfig": {
    "service": {
      "name": "machine-api-operator-webhook",
      "namespace": "openshift-machine-api",
      "path": "/validate-machine-openshift-io-v1beta1-machineset",
      "port": 443
    }
  },
  "failurePolicy": "Ignore",
  "matchPolicy": "Equivalent",
  "name": "validation.machineset.machine.openshift.io",
  "namespaceSelector": {},
  "objectSelector": {},
  "rules": [
    {
      "apiGroups": [
        "machine.openshift.io"
      ],
      "apiVersions": [
        "v1beta1"
      ],
      "operations": [
        "CREATE",
        "UPDATE"
      ],
      "resources": [
        "machinesets"
      ],
      "scope": "*"
    }
  ]
}
```

```
}
],
"sideEffects": "None",
"timeoutSeconds": 10
}
{
  "admissionReviewVersions": [
    "v1"
  ],
  "clientConfig": {
    "service": {
      "name": "multus-admission-controller",
      "namespace": "openshift-multus",
      "path": "/validate",
      "port": 443
    }
  },
  "failurePolicy": "Fail",
  "matchPolicy": "Equivalent",
  "name": "multus-validating-config.k8s.io",
  "namespaceSelector": {},
  "objectSelector": {},
  "rules": [
    {
      "apiGroups": [
        "k8s.cni.cncf.io"
      ],
      "apiVersions": [
        "v1"
      ],
      "operations": [
        "CREATE",
        "UPDATE"
      ],
      "resources": [
        "network-attachment-definitions"
      ],
      "scope": "*"
    }
  ],
  "sideEffects": "NoneOnDryRun",
  "timeoutSeconds": 30
}
{
  "admissionReviewVersions": [
    "v1"
  ],
  "clientConfig": {
    "service": {
      "name": "prometheus-operator",
      "namespace": "openshift-monitoring",
      "path": "/admission-prometheusrules/validate",
      "port": 8080
    }
  },
  "failurePolicy": "Ignore",
  "matchPolicy": "Equivalent",
```

```
"name": "prometheusrules.openshift.io",
"namespaceSelector": {},
"objectSelector": {},
"rules": [
  {
    "apiGroups": [
      "monitoring.coreos.com"
    ],
    "apiVersions": [
      "v1"
    ],
    "operations": [
      "CREATE",
      "UPDATE"
    ],
    "resources": [
      "prometheusrules"
    ],
    "scope": "Namespaced"
  }
],
"sideEffects": "None",
"timeoutSeconds": 5
}
```

## 7.12. Mutating WebHook Configurations

```
{
  "admissionReviewVersions": [
    "v1"
  ],
  "clientConfig": {
    "service": {
      "name": "machine-api-operator-webhook",
      "namespace": "openshift-machine-api",
      "path": "/mutate-machine-openshift-io-v1beta1-machine",
      "port": 443
    }
  },
  "failurePolicy": "Ignore",
  "matchPolicy": "Equivalent",
  "name": "default.machine.machine.openshift.io",
  "namespaceSelector": {},
  "objectSelector": {},
  "reinvocationPolicy": "Never",
  "rules": [
    {
      "apiGroups": [
        "machine.openshift.io"
      ],
      "apiVersions": [
        "v1beta1"
      ],
      "operations": [
```

```
        "CREATE"
      ],
      "resources": [
        "machines"
      ],
      "scope": "*"
    }
  ],
  "sideEffects": "None",
  "timeoutSeconds": 10
}
{
  "admissionReviewVersions": [
    "v1"
  ],
  "clientConfig": {
    "service": {
      "name": "machine-api-operator-webhook",
      "namespace": "openshift-machine-api",
      "path": "/mutate-machine-openshift-io-v1beta1-machineset",
      "port": 443
    }
  },
  "failurePolicy": "Ignore",
  "matchPolicy": "Equivalent",
  "name": "default.machineset.machine.openshift.io",
  "namespaceSelector": {},
  "objectSelector": {},
  "reinvocationPolicy": "Never",
  "rules": [
    {
      "apiGroups": [
        "machine.openshift.io"
      ],
      "apiVersions": [
        "v1beta1"
      ],
      "operations": [
        "CREATE"
      ],
      "resources": [
        "machinesets"
      ],
      "scope": "*"
    }
  ],
  "sideEffects": "None",
  "timeoutSeconds": 10
}
```

## 7.13. API Versions

### 7.13.1. Verification of API files permissions

```
/etc/kubernetes/static-pod-resources/kube-apiserver-pod.yaml: [OK]  
/etc/kubernetes/static-pod-resources/kube-controller-manager-pod.yaml: [OK]  
/etc/kubernetes/static-pod-resources/kube-scheduler-pod.yaml: [OK]  
/etc/kubernetes/manifests/etcd-pod.yaml: [OK]
```

## Chapter 8. storage

### 8.1. Failed Persistent Volumes

### 8.2. Pending Persistent Volumes

### 8.3. Released Persistent Volumes

### 8.4. Lost Persistent Volumes Claims

### 8.5. Pending Persistent Volumes Claims

### 8.6. Storage Classes

### 8.7. Quotas

NAMESPACE	NAME	AGE	REQUEST
LIMIT			
openshift-host-network	host-network-namespace-quotas	41d	count/daemonsets.apps: 0/0, count/deployments.apps: 0/0, pods: 0/0 limits.cpu: 0/0, limits.memory: 0/0

### 8.8. Volume Snapshots



## 8.9. CSI Drivers

--

## 8.10. CSI Nodes

NAME	DRIVERS	AGE
crc-5nvrn-master-0	0	41d

## 8.11. Feature Gates

--

## 8.12. Horizontal Pod AutoScalers

--

# Chapter 9. logging

## 9.1. Logging Resources

# Chapter 10. Monitoring

## 10.1. Prometheus Status

prometheus-adapter-6dbdddc76b-bdr2v	1/1	Running	0	19h
prometheus-k8s-0	6/6	Running	0	19h
prometheus-operator-76f79c8d85-xn57s	2/2	Running	0	19h

## 10.2. Prometheus Rules

### 10.2.1. openshift-cloud-credential-operator cloud-credential-operator-alerts

```
"CloudCredentialOperatorTargetNamespaceMissing"
"One ore more CredentialsRequest CRs are asking to save credentials to a non-existent namespace."
"cco_credentials_requests_conditions{condition=\"MissingTargetNamespace\"} > 0"
{"severity":"warning"}
"CloudCredentialOperatorProvisioningFailed"
"One or more CredentialsRequest CRs are unable to be processed."
"cco_credentials_requests_conditions{condition=\"CredentialsProvisionFailure\"} > 0"
{"severity":"warning"}
"CloudCredentialOperatorDeprovisioningFailed"
"One or more CredentialsRequest CRs are unable to be deleted."
"cco_credentials_requests_conditions{condition=\"CredentialsDeprovisionFailure\"} > 0"
{"severity":"warning"}
"CloudCredentialOperatorInsufficientCloudCreds"
"Problem with the available platform credentials."
"cco_credentials_requests_conditions{condition=\"InsufficientCloudCreds\"} > 0"
{"severity":"warning"}
"CloudCredentialOperatorStaleCredentials"
"One or more CredentialsRequest CRs are stale and should be deleted."
"cco_credentials_requests_conditions{condition=\"StaleCredentials\"} > 0"
{"severity":"warning"}
```

### 10.2.2. openshift-cluster-machine-approver machineapprover-rules

```
"MachineApproverMaxPendingCSRReached"
"max pending CSRs threshold reached."
"mapi_current_pending_csr > mapi_max_pending_csr\n"
{"severity":"warning"}
```

### 10.2.3. openshift-cluster-node-tuning-operator node-tuning-operator

```
"NTOPodsNotReady"
"Pod {{ $labels.pod }} is not ready."
"kube_pod_status_ready{namespace='openshift-cluster-node-tuning-operator', condition='true'} == 0\n"
{"severity":"warning"}
"NTODegraded"
"The Node Tuning Operator is degraded."
"nto_degraded_info == 1"
{"severity":"warning"}
null
null
"count(nto_profile_calculated_total{profile!~\"openshift-node\",profile!~\"openshift-control-plane\",profile!~\"openshift\"})"
null
```

#### 10.2.4. openshift-cluster-samples-operator samples-operator-alerts

```
"SamplesRetriesMissingOnImagestreamImportFailing"
null
"sum(openshift_samples_failed_imagestream_import_info) >
sum(openshift_samples_retry_imagestream_import_total) -
sum(openshift_samples_retry_imagestream_import_total offset 30m)"
{"severity":"warning"}
"SamplesImagestreamImportFailing"
null
"sum(openshift_samples_retry_imagestream_import_total) -
sum(openshift_samples_retry_imagestream_import_total offset 30m) >
sum(openshift_samples_failed_imagestream_import_info)"
{"severity":"warning"}
"SamplesDegraded"
null
"openshift_samples_degraded_info == 1"
{"severity":"warning"}
"SamplesInvalidConfig"
null
"openshift_samples_invalidconfig_info == 1"
{"severity":"warning"}
"SamplesMissingSecret"
null
"openshift_samples_invalidsecret_info{reason=\"missing_secret\"} == 1"
{"severity":"warning"}
"SamplesMissingTBRCredential"
null
"openshift_samples_invalidsecret_info{reason=\"missing_tbr_credential\"} == 1"
{"severity":"warning"}
"SamplesTBRIInaccessibleOnBoot"
null
"openshift_samples_tbr_inaccessible_info == 1"
{"severity":"info"}
```

#### 10.2.5. openshift-cluster-version cluster-version-operator

```
"ClusterVersionOperatorDown"
"Cluster version operator has disappeared from Prometheus target discovery."
"absent(up{job=\"cluster-version-operator\"} == 1)\n"
{"severity":"critical"}
"CannotRetrieveUpdates"
"Cluster version operator has not retrieved updates in {{ $value | humanizeDuration }}."
"(time()-cluster_version_operator_update_retrieval_timestamp_seconds) >= 3600 and ignoring(condition,
name, reason) cluster_operator_conditions{name=\"version\", condition=\"RetrievedUpdates\",
endpoint=\"metrics\", reason!=\"NoChannel\"})\n"
{"severity":"warning"}
"UpdateAvailable"
"Your upstream update recommendation service recommends you update your cluster."
"sum by (channel,upstream) (cluster_version_available_updates) > 0\n"
{"severity":"info"}
"ClusterNotUpgradeable"
"One or more cluster operators have been blocking minor version cluster upgrades for at least an hour."
"max by (name, condition, endpoint) (cluster_operator_conditions{name=\"version\",
condition=\"Upgradeable\", endpoint=\"metrics\"} == 0)\n"
{"severity":"info"}
"ClusterOperatorDown"
"Cluster operator has not been available for 10 minutes."
"cluster_operator_up{job=\"cluster-version-operator\"} == 0\n"
{"severity":"critical"}
"ClusterOperatorDegraded"
"Cluster operator has been degraded for 30 minutes."
"(\n cluster_operator_conditions{job=\"cluster-version-operator\", condition=\"Degraded\"})\n or on
(name)\n group by (name) (cluster_operator_up{job=\"cluster-version-operator\"})\n) == 1\n"
{"severity":"warning"}
"ClusterOperatorFlapping"
"Cluster operator up status is changing often."
"changes(cluster_operator_up{job=\"cluster-version-operator\"}[2m]) > 2\n"
{"severity":"warning"}
```

## 10.2.6. openshift-dns-operator dns

```
"CoreDNSPanicking"
"CoreDNS panic"
"increase(coredns_panics_total[10m]) > 0"
{"severity":"warning"}
"CoreDNSHealthCheckSlow"
"CoreDNS health checks"
"histogram_quantile(.95, sum(rate(coredns_health_request_duration_seconds_bucket[5m])) by (instance, le))
> 10"
{"severity":"warning"}
"CoreDNSErrorsHigh"
"CoreDNS serverfail"
"(sum by(namespace) (rate(coredns_dns_responses_total{rcode=\"SERVFAIL\"}[5m])))\n /\nsum by(namespace)
(rate(coredns_dns_responses_total[5m])))\n> 0.01\n"
{"severity":"warning"}
```

## 10.2.7. openshift-etcd-operator etcd-prometheus-rules

```
"etcdMembersDown"
"etcd cluster members are down."
"max without (endpoint) (\n sum without (instance) (up{job=~\".*etcd.*\"} == bool 0)\nor\n count without
(To) (\n sum without (instance) (rate(etcd_network_peer_sent_failures_total{job=~\".*etcd.*\"}[120s]))
> 0.01\n )\n)\n> 0\n"
{"severity":"critical"}
"etcdNoLeader"
"etcd cluster has no leader."
"etcd_server_has_leader{job=~\".*etcd.*\"} == 0\n"
{"severity":"critical"}
"etcdGRPCRequestsSlow"
"etcd grpc requests are slow"
"histogram_quantile(0.99, sum(rate(grpc_server_handling_seconds_bucket{job=~\".*etcd.*\",
grpc_method!=\"Defragment\", grpc_type=\"unary\"}[5m])) without(grpc_type))\n> 0.15\n"
{"severity":"critical"}
"etcdMemberCommunicationSlow"
"etcd cluster member communication is slow."
"histogram_quantile(0.99,
rate(etcd_network_peer_round_trip_time_seconds_bucket{job=~\".*etcd.*\"}[5m]))\n> 0.15\n"
{"severity":"warning"}
"etcdHighNumberOfFailedProposals"
"etcd cluster has high number of proposal failures."
"rate(etcd_server_proposals_failed_total{job=~\".*etcd.*\"}[15m]) > 5\n"
{"severity":"warning"}
"etcdHighFsyncDurations"
"etcd cluster 99th percentile fsync durations are too high."
"histogram_quantile(0.99, rate(etcd_disk_wal_fsync_duration_seconds_bucket{job=~\".*etcd.*\"}[5m]))\n>
0.5\n"
{"severity":"warning"}
"etcdHighFsyncDurations"
null
"histogram_quantile(0.99, rate(etcd_disk_wal_fsync_duration_seconds_bucket{job=~\".*etcd.*\"}[5m]))\n>
1\n"
{"severity":"critical"}
"etcdHighCommitDurations"
"etcd cluster 99th percentile commit durations are too high."
"histogram_quantile(0.99,
rate(etcd_disk_backend_commit_duration_seconds_bucket{job=~\".*etcd.*\"}[5m]))\n> 0.25\n"
{"severity":"warning"}
"etcdBackendQuotaLowSpace"
"etcd cluster database size exceeds the defined quota."
"(etcd_mvcc_db_total_size_in_bytes/etcd_server_quota_backend_bytes)*100 > 95\n"
{"severity":"critical"}
"etcdExcessiveDatabaseGrowth"
"etcd cluster database size increased by 50 percent over the past four hours."
"increase(((etcd_mvcc_db_total_size_in_bytes/etcd_server_quota_backend_bytes)*100)[240m:1m]) > 50\n"
{"severity":"warning"}
"etcdHighNumberOfFailedGRPCRequests"
"etcd cluster has high number of failed grpc requests."
"100 * sum(rate(grpc_server_handled_total{job=\"etcd\",
grpc_code=~\"Unknown|FailedPrecondition|ResourceExhausted|Internal|Unavailable|DataLoss|DeadlineExceeded\"
}[5m])) without (grpc_type, grpc_code)\n /\nsum(rate(grpc_server_handled_total{job=\"etcd\"}[5m]))
without (grpc_type, grpc_code)\n > 2 and on
()(sum(cluster_infrastructure_provider{type!~\"ipi|BareMetal\"} == bool 1))\n"
```

```
{ "severity": "warning" }
"etcdHighNumberOfFailedGRPCRequests"
"etcd cluster has high number of failed grpc requests."
"100 * sum(rate(grpc_server_handled_total{job=\"etcd\",
grpc_code=~\"Unknown|FailedPrecondition|ResourceExhausted|Internal|Unavailable|DataLoss|DeadlineExceeded\"
}[5m])) without (grpc_type, grpc_code)\n /\nsum(rate(grpc_server_handled_total{job=\"etcd\"}[5m]))
without (grpc_type, grpc_code)\n > 5 and on
() (sum(cluster_infrastructure_provider{type!~\"ipi|BareMetal\"} == bool 1))\n"
{ "severity": "critical" }
"etcdHighNumberOfLeaderChanges"
"etcd cluster has high number of leader changes."
"increase((max without (instance) (etcd_server_leader_changes_seen_total{job=~\".*etcd.*\"}) or
0*absent(etcd_server_leader_changes_seen_total{job=~\".*etcd.*\"}))[15m:1m]) >= 5\n"
{ "severity": "warning" }
"etcdInsufficientMembers"
"etcd is reporting that a majority of instances are unavailable."
"sum(up{job=\"etcd\"} == bool 1 and etcd_server_has_leader{job=\"etcd\"} == bool 1) without (instance,pod)
< ((count(up{job=\"etcd\"}) without (instance,pod) + 1) / 2)"
{ "severity": "critical" }
```

### 10.2.8. openshift-image-registry image-registry-operator-alerts

```
"ImageRegistryStorageReconfigured"
null
"increase(image_registry_operator_storage_reconfigured_total[30m]) > 0"
{ "severity": "warning" }
```

### 10.2.9. openshift-ingress-operator ingress-operator

```
"HAProxyReloadFail"
"HAProxy reload failure"
"template_router_reload_failure == 1"
{ "severity": "warning" }
"HAProxyDown"
"HAProxy is down"
"haproxy_up == 0"
{ "severity": "critical" }
"IngressControllerDegraded"
"IngressController is degraded"
"ingress_controller_conditions{condition=\"Degraded\"} == 1"
{ "severity": "warning" }
"IngressControllerUnavailable"
"IngressController is unavailable"
"ingress_controller_conditions{condition=\"Available\"} == 0"
{ "severity": "warning" }
```

### 10.2.10. openshift-insights insights-prometheus-rules

```
"InsightsDisabled"
"Insights operator is disabled."
"cluster_operator_conditions{name=\"insights\", condition=\"Disabled\"} == 1\n"
{"severity":"info"}
"SimpleContentAccessNotAvailable"
"Simple content access certificates are not available."
"max_over_time(cluster_operator_conditions{name=\"insights\", condition=\"SCANotAvailable\",
reason=\"NotFound\"}[5m]) == 1\n"
{"severity":"info"}
```

### 10.2.11. openshift-kube-apiserver-operator kube-apiserver-operator

```
"TechPreviewNoUpgrade"
"Cluster has enabled tech preview features that will prevent upgrades."
"cluster_feature_set{name!=\"\", namespace=\"openshift-kube-apiserver-operator\"} == 0\n"
{"severity":"warning"}
```

### 10.2.12. openshift-kube-apiserver api-usage

```
"APIRemovedInNextReleaseInUse"
"Deprecated API that will be removed in the next version is being used."
"group(apiserver_requested_deprecated_apis{removed_release=\"1.24\"}) by (group,version,resource) and (sum
by(group,version,resource) (rate(apiserver_request_total{system_client!=\"kube-controller-
manager\",system_client!=\"cluster-policy-controller\"}[4h]))) > 0\n"
{"namespace":"openshift-kube-apiserver","severity":"info"}
"APIRemovedInNextEUSReleaseInUse"
"Deprecated API that will be removed in the next EUS version is being used."
"group(apiserver_requested_deprecated_apis{removed_release=~\"1\\\\.2[45]\"}) by (group,version,resource)
and (sum by(group,version,resource) (rate(apiserver_request_total{system_client!=\"kube-controller-
manager\",system_client!=\"cluster-policy-controller\"}[4h]))) > 0\n"
{"namespace":"openshift-kube-apiserver","severity":"info"}
```

### 10.2.13. openshift-kube-apiserver audit-errors

```
"AuditLogError"
"An API Server instance was unable to write audit logs. This could be triggered by the node running out
of space, or a malicious actor tampering with the audit logs."
"sum by (apiserver,instance)(rate(apiserver_audit_error_total{apiserver=~\".+apiserver\"}[5m])) / sum by
(apiserver,instance) (rate(apiserver_audit_event_total{apiserver=~\".+apiserver\"}[5m])) > 0\n"
{"namespace":"openshift-kube-apiserver","severity":"warning"}
```

### 10.2.14. openshift-kube-apiserver cpu-utilization



```
"HighOverallControlPlaneCPU"
"CPU utilization across all three control plane nodes is higher than two control plane nodes can sustain;
a single control plane node outage may cause a cascading failure; increase available CPU."
"sum(\n 100 - (avg by (instance) (rate(node_cpu_seconds_total{mode=\"idle\"}[1m])) * 100)\n AND on
(instance) label_replace( kube_node_role{role=\"master\"}, \"instance\", \"$1\", \"node\", \"(.+)\")
)\n)\n\ncount(kube_node_role{role=\"master\"})\n> 60\n"
{"namespace":"openshift-kube-apiserver","severity":"warning"}
"ExtremelyHighIndividualControlPlaneCPU"
"CPU utilization on a single control plane node is very high, more CPU pressure is likely to cause a
failover; increase available CPU."
"100 - (avg by (instance) (rate(node_cpu_seconds_total{mode=\"idle\"}[1m])) * 100) > 90 AND on (instance)
label_replace( kube_node_role{role=\"master\"}, \"instance\", \"$1\", \"node\", \"(.+)\") \n"
{"namespace":"openshift-kube-apiserver","severity":"warning"}
"ExtremelyHighIndividualControlPlaneCPU"
"Sustained high CPU utilization on a single control plane node, more CPU pressure is likely to cause a
failover; increase available CPU."
"100 - (avg by (instance) (rate(node_cpu_seconds_total{mode=\"idle\"}[1m])) * 100) > 90 AND on (instance)
label_replace( kube_node_role{role=\"master\"}, \"instance\", \"$1\", \"node\", \"(.+)\") \n"
{"namespace":"openshift-kube-apiserver","severity":"critical"}
```

## 10.2.15. openshift-kube-apiserver kube-apiserver-requests

```
null
null
"max_over_time(sum(apiserver_current_inflight_requests{apiserver=~\"openshift-apiserver|kube-apiserver\"})
by (apiserver,requestKind)[2m:])\n"
null
```

## 10.2.16. openshift-kube-apiserver kube-apiserver-slos-basic

```
"KubeAPIErrorBudgetBurn"
"The API server is burning too much error budget."
"sum(apiserver_request:burnrate1h) > (14.40 * 0.01000)\nand\nsum(apiserver_request:burnrate5m) > (14.40 *
0.01000)\n"
{"long":"1h","namespace":"openshift-kube-apiserver","severity":"critical","short":"5m"}
"KubeAPIErrorBudgetBurn"
"The API server is burning too much error budget."
"sum(apiserver_request:burnrate6h) > (6.00 * 0.01000)\nand\nsum(apiserver_request:burnrate30m) > (6.00 *
0.01000)\n"
{"long":"6h","namespace":"openshift-kube-apiserver","severity":"critical","short":"30m"}
null
null
"# error\nlabel_replace(\n
sum(rate(apiserver_request_total{job=\"apiserver\",verb=~\"LIST|GET\",code=~\"5..\"}[5m]))\n/
scalar(sum(rate(apiserver_request_total{job=\"apiserver\",verb=~\"LIST|GET\"}[5m])))\n, \"type\",
\"error\", \"_none\", \"\")\nnor\n# resource-scoped latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job=\"apiserver\",verb=~\"LIST|GET\",subresource!~\"proxy|log|exec\",scope=\"resource\"}[5m])))\n -\n
(sum(rate(apiserver_request_duration_seconds_bucket{job=\"apiserver\",verb=~\"LIST|GET\",subresource!~\"proxy|log|exec\",scope=\"resource\",le=\"1\"}[5m])) or vector(0))\n ) /
scalar(sum(rate(apiserver_request_total{job=\"apiserver\",verb=~\"LIST|GET\",subresource!~\"proxy|log|exec\"}[5m])))\n, \"type\", \"slow-resource\", \"_none\", \"\")\nnor\n# namespace-scoped
```

```

latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace"}[5m]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace",le="5"}[5m]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[5m]))\n, \type, \slow-namespace, \_none, \)\n\nor\n# cluster-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",scope="cluster"}[5m]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",scope="cluster",le="30"}[5m]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[5m]))\n, \type, \slow-cluster, \_none, \)\n\n"
{"verb":"read"}
null
null
"# error\nlabel_replace(\n
sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",code=~"5.."}[30m]))\n/
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[30m]))\n, \type,
"error", \_none, \)\n\nor\n# resource-scoped latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="resource"}[30m]))\n -\n
(sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="resource",le="1"}[30m])) or vector(0))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[30m]))\n, \type, \slow-resource, \_none, \)\n\nor\n# namespace-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace"}[30m]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace",le="5"}[30m]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[30m]))\n, \type, \slow-namespace, \_none, \)\n\nor\n# cluster-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",scope="cluster"}[30m]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",scope="cluster",le="30"}[30m]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[30m]))\n, \type, \slow-cluster, \_none, \)\n\n"
{"verb":"read"}
null
null
"# error\nlabel_replace(\n
sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",code=~"5.."}[1h]))\n/
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[1h]))\n, \type,
"error", \_none, \)\n\nor\n# resource-scoped latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="resource"}[1h]))\n -\n
(sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="resource",le="1"}[1h])) or vector(0))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[1h]))\n, \type, \slow-resource, \_none, \)\n\nor\n# namespace-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace"}[1h]))\n -

```

```

sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace",le="5s"}[1h]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[1h]))\n, \"type\", \"slow-namespace\", \"_none_\", \"\")\nor\n# cluster-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",scope="cluster"}[1h]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",scope="cluster",le="30s"}[1h]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[1h]))\n, \"type\", \"slow-cluster\", \"_none_\", \"\")\n"
{"verb":"read"}
null
null
"# error\nlabel_replace(\n
sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",code=~"5.."}[6h]))\n/
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[6h]))\n, \"type\",
\"error\", \"_none_\", \"\")\nor\n# resource-scoped latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="resource"}[6h]))\n -\n
(sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="resource",le="1s"}[6h])) or vector(0))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[6h]))\n, \"type\", \"slow-resource\", \"_none_\", \"\")\nor\n# namespace-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace"}[6h]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec",scope="namespace",le="5s"}[6h]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET",subresource!~"proxy|log|exec"}[6h]))\n, \"type\", \"slow-namespace\", \"_none_\", \"\")\nor\n# cluster-scoped
latency\nlabel_replace(\n (\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"LIST|GET",scope="cluster"}[6h]))\n -
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"LIST|GET",scope="cluster",le="30s"}[6h]))\n ) /
scalar(sum(rate(apiserver_request_total{job="apiserver",verb=~"LIST|GET"}[6h]))\n, \"type\", \"slow-cluster\", \"_none_\", \"\")\n"
{"verb":"read"}
null
null
"(\n (\n # too slow\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"POST|PUT|PATCH|DELETE"}[1h]))\n
\n -\n
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"POST|PUT|PATCH|DELETE",le="1s"}[1h]))\n )\n +\n
sum(rate(apiserver_request_total{job="apiserver",verb=~"POST|PUT|PATCH|DELETE",code=~"5.."}[1h]))\n)\n
\nsum(rate(apiserver_request_total{job="apiserver",verb=~"POST|PUT|PATCH|DELETE"}[1h]))\n"
{"verb":"write"}
null
null
"(\n (\n # too slow\n
sum(rate(apiserver_request_duration_seconds_count{job="apiserver",verb=~"POST|PUT|PATCH|DELETE"}[30m]))\n
)\n -\n
sum(rate(apiserver_request_duration_seconds_bucket{job="apiserver",verb=~"POST|PUT|PATCH|DELETE",le="1s"}[30m]))\n )\n +\n
sum(rate(apiserver_request_total{job="apiserver",verb=~"POST|PUT|PATCH|DELETE",code=~"5.."}[30m]))\n)\n

```

```

)\n/\nsum(rate/apiserver_request_total{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[30m]))\n"
{"verb":"write"}
null
null
"(\n (\n # too slow\n
sum(rate/apiserver_request_duration_seconds_count{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[5m]))
\n -\n
sum(rate/apiserver_request_duration_seconds_bucket{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE",le="\
1"}[5m]))\n )\n +\n
sum(rate/apiserver_request_total{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE",code=~"5.."}[5m]))\n)
\n/\nsum(rate/apiserver_request_total{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[5m]))\n"
{"verb":"write"}
null
null
"(\n (\n # too slow\n
sum(rate/apiserver_request_duration_seconds_count{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[6h]))
\n -\n
sum(rate/apiserver_request_duration_seconds_bucket{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE",le="\
1"}[6h]))\n )\n +\n
sum(rate/apiserver_request_total{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE",code=~"5.."}[6h]))\n)
\n/\nsum(rate/apiserver_request_total{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[6h]))\n"
{"verb":"write"}
null
null
"sum by (code,resource) (rate/apiserver_request_total{job="\apiserver",verb=~"LIST|GET"}[5m]))\n"
{"verb":"read"}
null
null
"sum by (code,resource)
(rate/apiserver_request_total{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[5m]))\n"
{"verb":"write"}
null
null
"histogram_quantile(0.99, sum by (le, resource)
(rate/apiserver_request_duration_seconds_bucket{job="\apiserver",verb=~"LIST|GET"}[5m])) > 0\n"
{"quantile":"0.99","verb":"read"}
null
null
"histogram_quantile(0.99, sum by (le, resource)
(rate/apiserver_request_duration_seconds_bucket{job="\apiserver",verb=~"POST|PUT|PATCH|DELETE"}[5m]))
> 0\n"
{"quantile":"0.99","verb":"write"}
null
null
"histogram_quantile(0.99,
sum(rate/apiserver_request_duration_seconds_bucket{job="\apiserver",subresource!="log",verb!~"LIST|WAT
CH|WATCHLIST|DELETERECOLLECTION|PROXY|CONNECT"}[5m])) without(instance, pod))\n"
{"quantile":"0.99"}
null
null
"histogram_quantile(0.9,
sum(rate/apiserver_request_duration_seconds_bucket{job="\apiserver",subresource!="log",verb!~"LIST|WAT
CH|WATCHLIST|DELETERECOLLECTION|PROXY|CONNECT"}[5m])) without(instance, pod))\n"
{"quantile":"0.9"}
null
null
"histogram_quantile(0.5,

```

```
sum(rate(apiserver_request_duration_seconds_bucket{job=\"apiserver\",subresource!=\"log\",verb!~\"LIST|WATCH|WATCHLIST|DELETECOLLECTION|PROXY|CONNECT\"}[5m])) without(instance, pod))\n\"{\n  \"quantile\": \"0.5\"\n}
```

### 10.2.17. openshift-kube-controller-manager-operator kube-controller-manager-operator

```
"KubeControllerManagerDown"
"Target disappeared from Prometheus target discovery."
"absent(up{job=\"kube-controller-manager\"} == 1)\n"
{"namespace": "openshift-kube-controller-manager", "severity": "critical"}
"PodDisruptionBudgetAtLimit"
"The pod disruption budget is preventing further disruption to pods."
"max by(namespace, poddisruptionbudget) (kube_poddisruptionbudget_status_current_healthy ==
kube_poddisruptionbudget_status_desired_healthy)\n"
{"severity": "warning"}
"PodDisruptionBudgetLimit"
"The pod disruption budget registers insufficient amount of pods."
"max by (namespace, poddisruptionbudget) (kube_poddisruptionbudget_status_current_healthy <
kube_poddisruptionbudget_status_desired_healthy)\n"
{"severity": "critical"}
```

### 10.2.18. openshift-kube-scheduler-operator kube-scheduler-operator

```
"KubeSchedulerDown"
"Target disappeared from Prometheus target discovery."
"absent(up{job=\"scheduler\"} == 1)\n"
{"namespace": "openshift-kube-scheduler", "severity": "critical"}
"SchedulerLegacyPolicySet"
"Legacy scheduler policy API in use by the scheduler."
"cluster_legacy_scheduler_policy > 0\n"
{"severity": "warning"}
```

### 10.2.19. openshift-machine-api machine-api-operator-prometheus-rules

```
"MachineWithoutValidNode"
"machine {{ $labels.name }} does not have valid node reference"
"(mapi_machine_created_timestamp_seconds unless on(node) kube_node_info) > 0\n"
{"severity":"warning"}
"MachineWithNoRunningPhase"
"machine {{ $labels.name }} is in phase: {{ $labels.phase }}"
"(mapi_machine_created_timestamp_seconds{phase!~\"Running|Deleting\"}) > 0\n"
{"severity":"warning"}
"MachineNotYetDeleted"
"machine {{ $labels.name }} has been in Deleting phase for more than 6 hours"
"sum by (name, namespace) (avg_over_time(mapi_machine_created_timestamp_seconds{phase=\"Deleting\"}[15m]))
> 0\n"
{"severity":"warning"}
"MachineAPIOperatorMetricsCollectionFailing"
"machine api operator metrics collection is failing."
"mapi_mao_collector_up == 0\n"
{"severity":"critical"}
"MachineHealthCheckUnterminatedShortCircuit"
"machine health check {{ $labels.name }} has been disabled by short circuit for more than 30 minutes"
"mapi_machinehealthcheck_short_circuit == 1\n"
{"severity":"warning"}
```

## 10.2.20. openshift-machine-config-operator machine-config-daemon

```
"MCDRebootError"
null
"mcd_reboot_err > 0\n"
{"severity":"critical"}
"MCDDrainError"
null
"mcd_drain_err > 0\n"
{"severity":"warning"}
"MCDPivotError"
null
"mcd_pivot_err > 0\n"
{"severity":"warning"}
"KubeletHealthState"
null
"mcd_kubelet_state > 2\n"
{"severity":"warning"}
"SystemMemoryExceedsReservation"
null
"sum by (node) (container_memory_rss{id=\"/system.slice\"}) > ((sum by (node)
(kube_node_status_capacity{resource=\"memory\"} - kube_node_status_allocatable{resource=\"memory\"})) *
0.95)\n"
{"severity":"warning"}
"MasterNodesHighMemoryUsage"
null
"((sum(node_memory_MemTotal_bytes AND on (instance) label_replace( kube_node_role{role=\"master\"},
\"instance\", \"$1\", \"node\", \"(.+)\") ) - sum(node_memory_MemFree_bytes + node_memory_Buffers_bytes +
node_memory_Cached_bytes AND on (instance) label_replace( kube_node_role{role=\"master\"}, \"instance\",
\"$1\", \"node\", \"(.+)\") )) / sum(node_memory_MemTotal_bytes AND on (instance) label_replace(
kube_node_role{role=\"master\"}, \"instance\", \"$1\", \"node\", \"(.+)\") ) * 100) > 90\n"
{"severity":"warning"}
```

## 10.2.21. openshift-marketplace marketplace-alert-rules

```
"CommunityOperatorsCatalogError"
"Community-operators not ready for 10 minutes"
"catalogsource_ready{name=\"community-operators\",exported_namespace=\"openshift-marketplace\"} == 0"
{"severity":"warning"}
"CertifiedOperatorsCatalogError"
"Certified-operators not ready for more than 10 minutes"
"catalogsource_ready{name=\"certified-operators\",exported_namespace=\"openshift-marketplace\"} == 0"
{"severity":"warning"}
"RedhatOperatorsCatalogError"
"Redhat-operators not ready for more than 10 minutes"
"catalogsource_ready{name=\"redhat-operators\",exported_namespace=\"openshift-marketplace\"} == 0"
{"severity":"warning"}
"RedhatMarketplaceCatalogError"
"Redhat-marketplace not ready for more than 10 minutes"
"catalogsource_ready{name=\"redhat-marketplace\",exported_namespace=\"openshift-marketplace\"} == 0"
{"severity":"warning"}
```

## 10.2.22. openshift-monitoring alertmanager-main-rules

```
"AlertmanagerFailedReload"
"Reloading an Alertmanager configuration has failed."
"# Without max_over_time, failed scrapes could create false negatives, see\n#
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for
details.\nmax_over_time(alertmanager_config_last_reload_successful{job=\"alertmanager-
main\",namespace=\"openshift-monitoring\"}[5m]) == 0\n"
{"severity":"critical"}
"AlertmanagerMembersInconsistent"
"A member of an Alertmanager cluster has not found all other cluster members."
"# Without max_over_time, failed scrapes could create false negatives, see\n#
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for details.\n
max_over_time(alertmanager_cluster_members{job=\"alertmanager-main\",namespace=\"openshift-
monitoring\"}[5m])\n< on (namespace,service) group_left\n count by (namespace,service)
(max_over_time(alertmanager_cluster_members{job=\"alertmanager-main\",namespace=\"openshift-
monitoring\"}[5m]))\n"
{"severity":"warning"}
"AlertmanagerFailedToSendAlerts"
"An Alertmanager instance failed to send notifications."
"(\n rate(alertmanager_notifications_failed_total{job=\"alertmanager-main\",namespace=\"openshift-
monitoring\"}[5m])\n/\n rate(alertmanager_notifications_total{job=\"alertmanager-
main\",namespace=\"openshift-monitoring\"}[5m])\n)\n> 0.01\n"
{"severity":"warning"}
"AlertmanagerClusterFailedToSendAlerts"
"All Alertmanager instances in a cluster failed to send notifications to a critical integration."
"min by (namespace,service, integration) (\n
rate(alertmanager_notifications_failed_total{job=\"alertmanager-main\",namespace=\"openshift-monitoring\",
integration=~\".*\"}[5m])\n/\n rate(alertmanager_notifications_total{job=\"alertmanager-
main\",namespace=\"openshift-monitoring\", integration=~\".*\"}[5m])\n)\n> 0.01\n"
{"severity":"warning"}
"AlertmanagerConfigInconsistent"
"Alertmanager instances within the same cluster have different configurations."
"count by (namespace,service) (\n count_values by (namespace,service) (\n\"config_hash\",
alertmanager_config_hash{job=\"alertmanager-main\",namespace=\"openshift-monitoring\"})\n)\n!= 1\n"
{"severity":"warning"}
"AlertmanagerClusterDown"
"Half or more of the Alertmanager instances within the same cluster are down."
"(\n count by (namespace,service) (\n avg_over_time(up{job=\"alertmanager-
main\",namespace=\"openshift-monitoring\"}[5m]) < 0.5\n )\n/\n count by (namespace,service) (\n
up{job=\"alertmanager-main\",namespace=\"openshift-monitoring\"}\n )\n)\n>= 0.5\n"
{"severity":"warning"}
```

## 10.2.23. openshift-monitoring cluster-monitoring-operator-prometheus-rules

```
"TargetDown"
"Some targets were not reachable from the monitoring server for an extended period of time."
"100 * (count(up == 0 unless on (node) max by (node) (kube_node_spec_unschedulable == 1)) BY (job,
namespace, service) /\n count(up unless on (node) max by (node) (kube_node_spec_unschedulable == 1)) BY
(job, namespace, service)) > 10\n"
{"severity":"warning"}
null
```



```

null
"sum(rate(container_cpu_usage_seconds_total{container!="",pod!="\"}[5m])) BY (pod, namespace)"
null
null
null
"sum(container_fs_usage_bytes{pod!="\"}) BY (pod, namespace)"
null
null
null
"sum(container_memory_usage_bytes{container!="\"}) BY (namespace)"
null
null
null
"sum(rate(container_cpu_usage_seconds_total{container!="POD\",container!="\"}[5m])) BY (namespace)"
null
null
null
"sum(container_memory_usage_bytes{container=\"\",pod!=\"\"}) BY (cluster) / sum(machine_memory_bytes) BY (cluster)"
null
null
null
"sum(container_spec_cpu_shares{container=\"\",pod!=\"\"}) / 1000 / sum(machine_cpu_cores)"
null
null
null
"sum(rate(container_cpu_usage_seconds_total{container=\"\",pod!=\"\"}[5m])) / sum(machine_cpu_cores)"
null
null
null
"max without(endpoint, instance, job, pod, service) (kube_node_labels and on(node)
kube_node_role{role=\"master\"})"
{"label_node_role_kubernetes_io":"master","label_node_role_kubernetes_io_master":"true"}
null
null
"max without(endpoint, instance, job, pod, service) (kube_node_labels and on(node)
kube_node_role{role=\"infra\"})"
{"label_node_role_kubernetes_io_infra":"true"}
null
null
"max without(endpoint, instance, job, pod, service) (cluster:master_nodes and on(node)
cluster:infra_nodes)"
{"label_node_role_kubernetes_io_infra":"true","label_node_role_kubernetes_io_master":"true"}
null
null
"cluster:master_infra_nodes or on (node) cluster:master_nodes or on (node) cluster:infra_nodes or on
(node) max without(endpoint, instance, job, pod, service) (kube_node_labels)"
null
null
null
"kube_node_labels and on(node) (sum(label_replace(node_cpu_info, \"node\\\", \"$1\\\", \"instance\\\",
\"(.*)\\\")) by (node, package, core) == 2)"
{"label_node_hyperv_thread_enabled":"true"}
null
null
"count(sum(virt_platform) by (instance, type, system_manufacturer, system_product_name,
baseboard_manufacturer, baseboard_product_name)) by (type, system_manufacturer, system_product_name,

```

```

baseboard_manufacturer, baseboard_product_name)"
null
null
null
"sum by(label_beta_kubernetes_io_instance_type, label_node_role_kubernetes_io, label_kubernetes_io_arch,
label_node_openshift_io_os_id) (\n (\n cluster:master_nodes\n * on(node) group_left() max
by(node)\n (\n kube_node_status_capacity{resource=\"cpu\",unit=\"core\"}\n )\n )\n or
on(node) (\n max without(endpoint, instance, job, pod, service)\n (\n kube_node_labels\n ) *
on(node) group_left() max by(node)\n (\n
kube_node_status_capacity{resource=\"cpu\",unit=\"core\"}\n )\n )\n)\n"
null
null
null
"clamp_max(\n label_replace(\n sum by(instance, package, core) (\n
node_cpu_info{core!=\"\",package!=\"\"}\n or\n # Assume core = cpu and package = 0 for platforms
that don't expose core/package labels.\n
label_replace(label_join(node_cpu_info{core=\"\",package=\"\"}, \"core\", \"\", \"cpu\"), \"package\",
\"0\", \"package\", \"\")\n ) > 1,\n \"label_node_hyperthread_enabled\", \n \"true\", \n
\"instance\", \n \"(.*)\" \n ) or on (instance, package)\n label_replace(\n sum by(instance,
package, core) (\n label_replace(node_cpu_info{core!=\"\",package!=\"\"}\n or\n # Assume
core = cpu and package = 0 for platforms that don't expose core/package labels.\n
label_join(node_cpu_info{core=\"\",package=\"\"}, \"core\", \"\", \"cpu\"), \"package\", \"0\",
\"package\", \"\")\n ) <= 1,\n \"label_node_hyperthread_enabled\", \n \"false\", \n
\"instance\", \n \"(.*)\" \n ), \n 1\n)\n"
null
null
null
"topk by(node) (1, cluster:nodes_roles) * on (node)\n group_right(
label_beta_kubernetes_io_instance_type, label_node_role_kubernetes_io, label_node_openshift_io_os_id,
label_kubernetes_io_arch,\n label_node_role_kubernetes_io_master,
label_node_role_kubernetes_io_infra)\nlabel_replace( cluster:cpu_core_hyperthreading, \"node\", \"$1\",
\"instance\", \"(.*)\" )\n"
null
null
null
"count(cluster:cpu_core_node_labels) by (label_beta_kubernetes_io_instance_type,
label_node_hyperthread_enabled)"
null
null
null
"sum by(label_beta_kubernetes_io_instance_type, label_node_role_kubernetes_io)\n(\n (\n
cluster:master_nodes\n * on(node) group_left() max by(node)\n (\n
kube_node_status_capacity{resource=\"memory\",unit=\"byte\"}\n )\n )\n or on(node)\n (\n max
without(endpoint, instance, job, pod, service)\n (\n kube_node_labels\n )\n * on(node)
group_left() max by(node)\n (\n kube_node_status_capacity{resource=\"memory\",unit=\"byte\"}\n
)\n )\n)\n"
null
null
null
"sum(1 - rate(node_cpu_seconds_total{mode=\"idle\"}[2m]) * on(namespace, pod) group_left(node)
node_namespace_pod:kube_pod_info:{pod=~\"node-exporter.+\"})"
null
null
null
"sum(node_memory_MemTotal_bytes{job=\"node-exporter\"} - node_memory_MemAvailable_bytes{job=\"node-
exporter\"})"
null

```

```
null
null
"sum(rate(container_cpu_usage_seconds_total{namespace!~\"openshift-.+\",pod!=\"\",container=\\\"\\\"}[5m]))"
null
null
null
"cluster:cpu_usage_cores:sum - workload:cpu_usage_cores:sum"
null
null
null
"sum(container_memory_working_set_bytes{namespace!~\"openshift-.+\",pod!=\"\",container=\\\"\\\"})"
null
null
null
"cluster:memory_usage_bytes:sum - workload:memory_usage_bytes:sum"
null
null
null
"sum(cluster:master_nodes or on(node) kube_node_labels ) BY (label_beta_kubernetes_io_instance_type,
label_node_role_kubernetes_io, label_kubernetes_io_arch, label_node_openshift_io_os_id)"
null
null
null
"sum by(provisioner) (\n topk by (namespace, persistentvolumeclaim) (\n 1,
kube_persistentvolumeclaim_resource_requests_storage_bytes\n ) * on(namespace, persistentvolumeclaim)
group_right()\n topk by(namespace, persistentvolumeclaim) (\n 1, kube_persistentvolumeclaim_info *
on(storageclass) group_left(provisioner) topk by(storageclass) (1, max by(storageclass, provisioner)
(kube_storageclass_info))\n )\n)\n"
null
null
null
"(sum(node_role_os_version_machine:cpu_capacity_cores:sum{label_node_role_kubernetes_io_master=\\\"\\\",label_
node_role_kubernetes_io_infra=\\\"\\\"} or absent(__does_not_exist__)*0)) +
((sum(node_role_os_version_machine:cpu_capacity_cores:sum{label_node_role_kubernetes_io_master=\\\"true\\\"}
or absent(__does_not_exist__)*0) * ((max(cluster_master_schedulable == 1)*0+1) or
(absent(cluster_master_schedulable == 1)*0))))"
null
null
null
"min_over_time(workload:capacity_physical_cpu_cores:sum[5m:15s])"
null
null
null
"max_over_time(workload:capacity_physical_cpu_cores:sum[5m:15s])"
null
null
null
"sum by (provisioner) (\n topk by (namespace, persistentvolumeclaim) (\n 1,
kubetlet_volume_stats_used_bytes\n ) * on (namespace,persistentvolumeclaim) group_right()\n topk by
(namespace, persistentvolumeclaim) (\n 1, kube_persistentvolumeclaim_info * on(storageclass)
group_left(provisioner) topk by(storageclass) (1, max by(storageclass, provisioner)
(kube_storageclass_info))\n )\n)\n"
null
null
null
"sum by (instance) (apiserver_storage_objects)"
null
```

```

null
null
"topk(500, max by(resource) (apiserver_storage_objects))"
null
null
null
"count(count (kube_pod_restart_policy{type!=\"Always\"},namespace!~\"openshift-.\+\"}) by (namespace,pod))"
null
null
null
"sum(max(kubelet_containers_per_pod_count_sum) by (instance))"
null
null
null
"count(cluster:cpu_core_node_labels) by (label_kubernetes_io_arch, label_node_hyperthread_enabled,
label_node_openshift_io_os_id,label_node_role_kubernetes_io_master,label_node_role_kubernetes_io_infra)"
null
null
null
"count(max(cluster:cpu_core_node_labels) by (node, package, label_beta_kubernetes_io_instance_type,
label_node_hyperthread_enabled, label_node_role_kubernetes_io) ) by (
label_beta_kubernetes_io_instance_type, label_node_hyperthread_enabled, label_node_role_kubernetes_io)"
null
null
null
"count (max(cluster:cpu_core_node_labels) by (node, package, label_kubernetes_io_arch,
label_node_hyperthread_enabled,
label_node_openshift_io_os_id,label_node_role_kubernetes_io_master,label_node_role_kubernetes_io_infra) )
by (label_kubernetes_io_arch, label_node_hyperthread_enabled,
label_node_openshift_io_os_id,label_node_role_kubernetes_io_master,label_node_role_kubernetes_io_infra)"
null
null
null
"max(alertmanager_integrations{namespace=\"openshift-monitoring\"})"
null
null
null
"sum by(plugin_name, volume_mode)(pv_collector_total_pv_count)"
null
null
null
"sum by(version)(vsphere_vcenter_info)"
null
null
null
"sum by(version)(vsphere_esxi_version_total)"
null
null
null
"sum by(hw_version)(vsphere_node_hw_version_total)"
null
null
null
"sum(\n min by (node) (kube_node_status_condition{condition=\"Ready\",status=\"true\"})\n and\n max
by (node) (kube_node_role{role=\"master\"})\n) == bool sum(kube_node_role{role=\"master\"})\n"
null
"ClusterMonitoringOperatorReconciliationErrors"

```

```

"Cluster Monitoring Operator is experiencing unexpected reconciliation errors."
"max_over_time(cluster_monitoring_operator_last_reconciliation_successful[5m]) == 0"
{"severity":"warning"}
"AlertmanagerReceiversNotConfigured"
"Receivers (notification integrations) are not configured on Alertmanager"
"cluster:alertmanager_integrations:max == 0"
{"namespace":"openshift-monitoring","severity":"warning"}
"KubeDeploymentReplicasMismatch"
"Deployment has not matched the expected number of replicas"
"(((\n kube_deployment_spec_replicas{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-state-
metrics\"}\n >\n kube_deployment_status_replicas_available{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kube-state-metrics\"}\n) and (\n
changes(kube_deployment_status_replicas_updated{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-
state-metrics\"}[5m])\n ==\n 0\n)) * on() group_left cluster:control_plane:all_nodes_ready) > 0\n"
{"severity":"warning"}
"MultipleContainersOOMKilled"
"Containers are being killed due to OOM"
"sum(max by(namespace, container, pod) (increase(kube_pod_container_status_restarts_total[12m])) and max
by(namespace, container, pod) (kube_pod_container_status_last_terminated_reason{reason=~\"OOMKilled\"}) ==
1) > 5"
{"namespace":"kube-system","severity":"info"}
null
null
"avg_over_time((((count((max by (node) (up{job=\"kubelet\",metrics_path=\"/metrics\"} == 1) and max by
(node) (kube_node_status_condition{condition=\"Ready\",status=\"true\"} == 1) and min by (node)
(kube_node_spec_unschedulable == 0))) / scalar(count(min by (node) (kube_node_spec_unschedulable ==
0))))))[5m:1s])"
null
null
null
"avg_over_time((count(max by (node) (kube_node_status_condition{condition=\"Ready\",status=\"true\"} ==
1)) / scalar(count(max by (node)
(kube_node_status_condition{condition=\"Ready\",status=\"true\"}))))[5m:1s])"
null
null
null
"(max without (condition,container,endpoint,instance,job,service)
(((kube_pod_status_ready{condition=\"false\"} == 1)*0 or (kube_pod_status_ready{condition=\"true\"} == 1))
* on(pod,namespace) group_left() group by (pod,namespace)
(kube_pod_status_phase{phase=~\"Running|Unknown|Pending\"} == 1)))"
null
null
null
"avg(kube_running_pod_ready{namespace=~\"openshift-.*\"})"
null
null
null
"avg(kube_running_pod_ready{namespace!~\"openshift-.*\"})"
null
null
null
"sum_over_time(workload:capacity_physical_cpu_cores:sum[30s:1s]) +
((cluster:usage:workload:capacity_physical_cpu_core_seconds offset 25s) or
(absent(cluster:usage:workload:capacity_physical_cpu_core_seconds offset 25s)*0))"
null
null
null

```

```

"sum by (code) (rate(haproxy_server_http_responses_total[5m]) > 0)"
null
null
null
"sum (rate(haproxy_frontend_bytes_in_total[5m]))"
null
null
null
"sum (rate(haproxy_frontend_bytes_out_total[5m]))"
null
null
null
"sum (haproxy_frontend_current_sessions)"
null
null
null
"sum(max without(service,endpoint,container,pod,job,namespace)
(increase(haproxy_server_http_responses_total{code!~\"2xx|1xx|4xx|3xx\",exported_namespace!~\"openshift-
.*\"}[5m]) > 0)) / sum (max without(service,endpoint,container,pod,job,namespace)
(increase(haproxy_server_http_responses_total{exported_namespace!~\"openshift-.*\"}[5m]))) or
absent(__does_not_exist__)*0"
null
null
null
"sum (max without(service,endpoint,container,pod,job,namespace)
(irate(haproxy_server_http_responses_total{exported_namespace!~\"openshift-.*\"}[5m]))) or
absent(__does_not_exist__)*0"
null
null
null
"sum(max without(service,endpoint,container,pod,job,namespace)
(increase(haproxy_server_http_responses_total{code!~\"2xx|1xx|4xx|3xx\",exported_namespace=~\"openshift-
.*\"}[5m]) > 0)) / sum (max without(service,endpoint,container,pod,job,namespace)
(increase(haproxy_server_http_responses_total{exported_namespace=~\"openshift-.*\"}[5m]))) or
absent(__does_not_exist__)*0"
null
null
null
"sum (max without(service,endpoint,container,pod,job,namespace)
(irate(haproxy_server_http_responses_total{exported_namespace=~\"openshift-.*\"}[5m]))) or
absent(__does_not_exist__)*0"
null
null
null
"sum(ingress_controller_aws_nlb_active) or vector(0)"
null
null
null
"sum by (strategy) (openshift_build_status_phase_total)"
null
null
null
"sum by (job,namespace) (max without(instance) (prometheus_tsdb_head_series{namespace=~\"openshift-
monitoring|openshift-user-workload-monitoring\"}))"
null
null
null

```

```
"sum by(job,namespace) (max without(instance)
(rate(prometheus_tsdb_head_samples_appended_total{namespace=~\"openshift-monitoring|openshift-user-
workload-monitoring\"}[2m])))"
null
null
null
"sum by (namespace) (max without(instance) (container_memory_working_set_bytes{namespace=~\"openshift-
monitoring|openshift-user-workload-monitoring\", container=~\"\"}))"
null
null
null
"topk(3, sum by(namespace, job)(sum_over_time(scrape_series_added[1h])))"
null
null
null
"topk(3, max by(namespace, job) (topk by(namespace,job) (1, scrape_samples_post_metric_relabeling)))"
null
null
null
"sum by(exported_service) (rate(haproxy_server_http_responses_total{exported_namespace=~\"openshift-
monitoring\", exported_service=~\"alertmanager-main|grafana|prometheus-k8s\"}[5m]))"
null
null
null
"max by (cluster, namespace, workload, pod) (label_replace(label_replace(kube_pod_owner{job=\"kube-state-
metrics\", owner_kind=\"ReplicationController\"},\"replicationcontroller\", \"$1\", \"owner_name\",
\"(.*)\") * on(replicationcontroller, namespace) group_left(owner_name) topk by(replicationcontroller,
namespace) (1, max by (replicationcontroller, namespace, owner_name)
(kube_replicationcontroller_owner{job=\"kube-state-metrics\"}),\"workload\", \"$1\", \"owner_name\",
\"(.*)\")))"
{"workload_type":"deploymentconfig"}
null
null
"sum by (instance) (etcd_mvcc_db_total_size_in_bytes{job=\"etcd\"})"
null
null
null
"histogram_quantile(0.99, sum by (instance, le)
(rate(etcd_disk_wal_fsync_duration_seconds_bucket{job=\"etcd\"}[5m])))"
{"quantile":"0.99"}
null
null
"histogram_quantile(0.99, sum by (instance, le)
(rate(etcd_network_peer_round_trip_time_seconds_bucket{job=\"etcd\"}[5m])))"
{"quantile":"0.99"}
null
null
"sum by (instance) (etcd_mvcc_db_total_size_in_use_in_bytes{job=\"etcd\"})"
null
null
null
"histogram_quantile(0.99, sum by (instance, le)
(rate(etcd_disk_backend_commit_duration_seconds_bucket{job=\"etcd\"}[5m])))"
{"quantile":"0.99"}
null
null
"sum(rate/apiserver_request_total{job=\"apiserver\"}[10m]) BY (code)"
```

```
null
"Watchdog"
"An alert that should always be firing to certify that Alertmanager is working properly."
"vector(1)"
{"namespace":"openshift-monitoring","severity":"none"}
"NodeNetworkInterfaceFlapping"
"Network interface is often changing its status"
"changes(node_network_up{job=\"node-exporter\",device!~\"veth.+\"}[2m]) > 2\n"
{"severity":"warning"}
null
null
"sum(rate(node_cpu_seconds_total{mode!=\"idle\",mode!=\"iowait\",mode!=\"steal\"}[3m])) BY (instance)"
null
null
null
"sum(rate(node_network_receive_bytes_total[3m])) BY (instance)"
null
null
null
"sum(rate(node_network_transmit_bytes_total[3m])) BY (instance)"
null
null
null
"sum(rate(node_cpu_seconds_total{mode!=\"idle\",mode!=\"iowait\",mode!=\"steal\"}[5m]))"
null
null
null
"cluster:node_cpu_seconds_total:rate5m / count(sum(node_cpu_seconds_total) BY (instance, cpu))"
null
null
null
"count without(instance, pod, node) (up == 1)"
null
null
null
"count without(instance, pod, node) (up == 0)"
null
```

## 10.2.24. openshift-monitoring kube-state-metrics-rules

```
"KubeStateMetricsListErrors"
"kube-state-metrics is experiencing errors in list operations."
"(sum(rate(kube_state_metrics_list_total{job=\"kube-state-metrics\",result=\"error\"}[5m])))\n
/\nsum(rate(kube_state_metrics_list_total{job=\"kube-state-metrics\"}[5m])))\n> 0.01\n"
{"severity":"warning"}
"KubeStateMetricsWatchErrors"
"kube-state-metrics is experiencing errors in watch operations."
"(sum(rate(kube_state_metrics_watch_total{job=\"kube-state-metrics\",result=\"error\"}[5m])))\n
/\nsum(rate(kube_state_metrics_watch_total{job=\"kube-state-metrics\"}[5m])))\n> 0.01\n"
{"severity":"warning"}
```



## 10.2.25. openshift-monitoring kubernetes-monitoring-rules

```
"KubePodCrashLooping"
"Pod is crash looping."
"max_over_time(kube_pod_container_status_waiting_reason{reason=\"CrashLoopBackOff\",
namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-metrics\"}[5m]) >= 1\n"
{"severity":"warning"}
"KubePodNotReady"
"Pod has been in a non-ready state for more than 15 minutes."
"sum by (namespace, pod) (\n max by(namespace, pod) (\n kube_pod_status_phase{namespace=~\"(openshift-
.*|kube-.*|default)\",job=\"kubernetes-metrics\", phase=~\"Pending|Unknown\"}\n ) * on(namespace, pod)
group_left(owner_kind) topk by(namespace, pod) (\n 1, max by(namespace, pod, owner_kind)
(kube_pod_owner{owner_kind!=\"Job\"})\n )\n) > 0\n"
{"severity":"warning"}
"KubeDeploymentGenerationMismatch"
"Deployment generation mismatch due to possible roll-back"
"kube_deployment_status_observed_generation{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n !=\nkube_deployment_metadata_generation{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n"
{"severity":"warning"}
"KubeStatefulSetReplicasMismatch"
"Deployment has not matched the expected number of replicas."
"(\n kube_statefulset_status_replicas_ready{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n !=\nkube_statefulset_status_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n) and (\n
changes(kube_statefulset_status_replicas_updated{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}[10m])\n ==\n 0\n)\n"
{"severity":"warning"}
"KubeStatefulSetGenerationMismatch"
"StatefulSet generation mismatch due to possible roll-back"
"kube_statefulset_status_observed_generation{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n !=\nkube_statefulset_metadata_generation{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n"
{"severity":"warning"}
"KubeStatefulSetUpdateNotRolledOut"
"StatefulSet update has not been rolled out."
"(\n max without (revision) (\n kube_statefulset_status_current_revision{namespace=~\"(openshift-
.*|kube-.*|default)\",job=\"kubernetes-metrics\"}\n unless\n
kube_statefulset_status_update_revision{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n )\n )\n * (\n kube_statefulset_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n !=\n
kube_statefulset_status_replicas_updated{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n )\n ) and (\n changes(kube_statefulset_status_replicas_updated{namespace=~\"(openshift-
.*|kube-.*|default)\",job=\"kubernetes-metrics\"}[5m])\n ==\n 0\n)\n"
{"severity":"warning"}
"KubeDaemonSetRolloutStuck"
"DaemonSet rollout is stuck."
"(\n (\n kube_daemonset_status_current_number_scheduled{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n !=\n
kube_daemonset_status_desired_number_scheduled{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n ) or (\n kube_daemonset_status_number_misscheduled{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n !=\n 0\n ) or (\n
kube_daemonset_updated_number_scheduled{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-
metrics\"}\n !=\n kube_daemonset_status_desired_number_scheduled{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubernetes-metrics\"}\n ) or (\n
kube_daemonset_status_number_available{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubernetes-"
```

```
metrics\"}\\n    !=\\n    kube_daemonset_status_desired_number_scheduled{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}\\n    )\\n) and (\\n
changes(kube_daemonset_updated_number_scheduled{namespace=~\"(openshift-.*|kube-.*|default)\\\",job=\\\"kube-
state-metrics\\\"}[5m])\\n    ==\\n    0\\n)\\n\"
{\"severity\":\"warning\"}
\"KubeContainerWaiting\"
\"Pod container waiting longer than 1 hour\"
\"sum by (namespace, pod, container) (kube_pod_container_status_waiting_reason{namespace=~\"(openshift-
.*|kube-.*|default)\\\",job=\\\"kube-state-metrics\\\"}) > 0\\n\"
{\"severity\":\"warning\"}
\"KubeDaemonSetNotScheduled\"
\"DaemonSet pods are not scheduled.\"
\"kube_daemonset_status_desired_number_scheduled{namespace=~\"(openshift-.*|kube-.*|default)\\\",job=\\\"kube-
state-metrics\\\"}\\n    -\\nkube_daemonset_status_current_number_scheduled{namespace=~\"(openshift-.*|kube
-.*|default)\\\",job=\\\"kube-state-metrics\\\"} > 0\\n\"
{\"severity\":\"warning\"}
\"KubeDaemonSetMisScheduled\"
\"DaemonSet pods are misscheduled.\"
\"kube_daemonset_status_number_misscheduled{namespace=~\"(openshift-.*|kube-.*|default)\\\",job=\\\"kube-state-
metrics\\\"} > 0\\n\"
{\"severity\":\"warning\"}
\"KubeJobCompletion\"
\"Job did not complete in time\"
\"kube_job_spec_completions{namespace=~\"(openshift-.*|kube-.*|default)\\\",job=\\\"kube-state-metrics\\\"} -
kube_job_status_succeeded{namespace=~\"(openshift-.*|kube-.*|default)\\\",job=\\\"kube-state-metrics\\\"} >
0\\n\"
{\"severity\":\"warning\"}
\"KubeJobFailed\"
\"Job failed to complete.\"
\"kube_job_failed{namespace=~\"(openshift-.*|kube-.*|default)\\\",job=\\\"kube-state-metrics\\\"} > 0\\n\"
{\"severity\":\"warning\"}
\"KubeHpaReplicasMismatch\"
\"HPA has not matched descired number of replicas.\"
\"(kube_horizontalpodautoscaler_status_desired_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}\\n
!=\\nkube_horizontalpodautoscaler_status_current_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"})\\n
and\\n(kube_horizontalpodautoscaler_status_current_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}\\n
>\\nkube_horizontalpodautoscaler_spec_min_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"})\\n
and\\n(kube_horizontalpodautoscaler_status_current_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}\\n
<\\nkube_horizontalpodautoscaler_spec_max_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"})\\n
and\\nchanges(kube_horizontalpodautoscaler_status_current_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}[15m]) == 0\\n\"
{\"severity\":\"warning\"}
\"KubeHpaMaxedOut\"
\"HPA is running at max replicas\"
\"kube_horizontalpodautoscaler_status_current_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}\\n
==\\nkube_horizontalpodautoscaler_spec_max_replicas{namespace=~\"(openshift-.*|kube-
.*|default)\\\",job=\\\"kube-state-metrics\\\"}\\n\"
{\"severity\":\"warning\"}
\"KubeCPUOvercommit\"
\"Cluster has overcommitted CPU resource requests.\"
```

```
"sum(namespace_cpu:kube_pod_container_resource_requests:sum{}) -
(sum(kube_node_status_allocatable{resource=\"cpu\"})) -
max(kube_node_status_allocatable{resource=\"cpu\"})) >
0\nand\n(sum(kube_node_status_allocatable{resource=\"cpu\"}) -
max(kube_node_status_allocatable{resource=\"cpu\"})) > 0\n"
{"namespace":"kube-system","severity":"warning"}
"KubeMemoryOvercommit"
"Cluster has overcommitted memory resource requests."
"sum(namespace_memory:kube_pod_container_resource_requests:sum{}) -
(sum(kube_node_status_allocatable{resource=\"memory\"})) -
max(kube_node_status_allocatable{resource=\"memory\"})) >
0\nand\n(sum(kube_node_status_allocatable{resource=\"memory\"}) -
max(kube_node_status_allocatable{resource=\"memory\"})) > 0\n"
{"namespace":"kube-system","severity":"warning"}
"KubeCPUQuotaOvercommit"
"Cluster has overcommitted CPU resource requests."
"sum(min without(resource) (kube_resourcequota{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-
state-metrics\", type=\"hard\", resource=~\"(cpu|requests.cpu)\"}))\n
/\nsum(kube_node_status_allocatable{resource=\"cpu\", job=\"kube-state-metrics\"})\n > 1.5\n"
{"severity":"warning"}
"KubeMemoryQuotaOvercommit"
"Cluster has overcommitted memory resource requests."
"sum(min without(resource) (kube_resourcequota{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-
state-metrics\", type=\"hard\", resource=~\"(memory|requests.memory)\"}))\n
/\nsum(kube_node_status_allocatable{resource=\"memory\", job=\"kube-state-metrics\"})\n > 1.5\n"
{"severity":"warning"}
"KubeQuotaAlmostFull"
"Namespace quota is going to be full."
"kube_resourcequota{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-state-metrics\",
type=\"used\"}\n / ignoring(instance, job, type)\n(kube_resourcequota{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kube-state-metrics\", type=\"hard\"} > 0)\n > 0.9 < 1\n"
{"severity":"info"}
"KubeQuotaFullyUsed"
"Namespace quota is fully used."
"kube_resourcequota{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-state-metrics\",
type=\"used\"}\n / ignoring(instance, job, type)\n(kube_resourcequota{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kube-state-metrics\", type=\"hard\"} > 0)\n == 1\n"
{"severity":"info"}
"KubeQuotaExceeded"
"Namespace quota has exceeded the limits."
"kube_resourcequota{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kube-state-metrics\",
type=\"used\"}\n / ignoring(instance, job, type)\n(kube_resourcequota{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kube-state-metrics\", type=\"hard\"} > 0)\n > 1\n"
{"severity":"warning"}
"KubePersistentVolumeFillingUp"
"PersistentVolume is filling up."
"(\n kubelet_volume_stats_available_bytes{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubelet\",
metrics_path=\"/metrics\"}\n /\n kubelet_volume_stats_capacity_bytes{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubelet\", metrics_path=\"/metrics\"}\n) <
0.03\nand\nkubelet_volume_stats_used_bytes{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubelet\",
metrics_path=\"/metrics\"} > 0\nunless on(namespace,
persistentvolumeclaim)\nkube_persistentvolumeclaim_access_mode{namespace=~\"(openshift-.*|kube-
.*|default)\", access_mode=\"ReadOnlyMany\"} == 1\nunless on(namespace,
persistentvolumeclaim)\nkube_persistentvolumeclaim_labels{namespace=~\"(openshift-.*|kube-
.*|default)\",label_alerts_k8s_io_kube_persistent_volume_filling_up=\"disabled\"} == 1\n"
{"severity":"critical"}
"KubePersistentVolumeFillingUp"
```

```

"PersistentVolume is filling up."
"(\n kubelet_volume_stats_available_bytes{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubelet\",
metrics_path=\"/metrics\"}\n / \n kubelet_volume_stats_capacity_bytes{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubelet\", metrics_path=\"/metrics\"}\n) <
0.15\nand\nkubelet_volume_stats_used_bytes{namespace=~\"(openshift-.*|kube-.*|default)\",job=\"kubelet\",
metrics_path=\"/metrics\"} >
0\nand\npredict_linear(kubelet_volume_stats_available_bytes{namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kubelet\", metrics_path=\"/metrics\"}[6h], 4 * 24 * 3600) < 0\nunless on(namespace,
persistentvolumeclaim)\nkube_persistentvolumeclaim_access_mode{namespace=~\"(openshift-.*|kube-
.*|default)\", access_mode=\"ReadOnlyMany\"} == 1\nunless on(namespace,
persistentvolumeclaim)\nkube_persistentvolumeclaim_labels{namespace=~\"(openshift-.*|kube-
.*|default)\",label_alerts_k8s_io_kube_persistent_volume_filling_up=\"disabled\"} == 1\n"
{"severity":"warning"}
"KubePersistentVolumeErrors"
"PersistentVolume is having issues with provisioning."
"kube_persistentvolume_status_phase{phase=~\"Failed|Pending\",namespace=~\"(openshift-.*|kube-
.*|default)\",job=\"kube-state-metrics\"} > 0\n"
{"severity":"warning"}
"KubeClientErrors"
"Kubernetes API server client is experiencing errors."
"(sum(rate(rest_client_requests_total{code=~\"5..\"}[5m])) by (instance, job, namespace))\n
/\nsum(rate(rest_client_requests_total[5m])) by (instance, job, namespace))\n> 0.01\n"
{"severity":"warning"}
"KubeAggregatedAPIErrors"
"Kubernetes aggregated API has reported errors."
"sum by(name, namespace)(increase(aggregator_unavailable_apiservice_total[10m])) > 4\n"
{"severity":"warning"}
"KubeAggregatedAPIDown"
"Kubernetes aggregated API is down."
"(1 - max by(name, namespace)(avg_over_time(aggregator_unavailable_apiservice[10m]))) * 100 < 85\n"
{"severity":"warning"}
"KubeAPIDown"
"Target disappeared from Prometheus target discovery."
"absent(up{job=\"apiserver\"} == 1)\n"
{"severity":"critical"}
"KubeAPITerminatedRequests"
"The kubernetes apiserver has terminated {{ $value | humanizePercentage }} of its incoming requests."
"sum(rate(apiserver_request_terminations_total{job=\"apiserver\"}[10m])) / (
sum(rate(apiserver_request_total{job=\"apiserver\"}[10m])) +
sum(rate(apiserver_request_terminations_total{job=\"apiserver\"}[10m])) ) > 0.20\n"
{"severity":"warning"}
"KubeNodeNotReady"
"Node is not ready."
"kube_node_status_condition{job=\"kube-state-metrics\",condition=\"Ready\",status=\"true\"} == 0\n"
{"severity":"warning"}
"KubeNodeUnreachable"
"Node is unreachable."
"(kube_node_spec_taint{job=\"kube-state-
metrics\",key=\"node.kubernetes.io/unreachable\",effect=\"NoSchedule\"} unless ignoring(key,value)
kube_node_spec_taint{job=\"kube-state-
metrics\",key=~\"ToBeDeletedByClusterAutoscaler|cloud.google.com/impending-node-termination|aws-node-
termination-handler/spot-itn\"}) == 1\n"
{"severity":"warning"}
"KubeletTooManyPods"
"Kubelet is running at capacity."
"count by(node) (\n kube_pod_status_phase{job=\"kube-state-metrics\",phase=\"Running\"} == 1) *
on(instance,pod,namespace,cluster) group_left(node) topk by(instance,pod,namespace,cluster) (1,

```

```
kube_pod_info{job=\"kube-state-metrics\\\")\\n\\n\\nmax by(node) (\\n kube_node_status_capacity{job=\"kube-
state-metrics\\\",resource=\"pods\\\"} != 1\\n) > 0.95\\n\"
{\"severity\":\"info\"}
\"KubeNodeReadinessFlapping\"
\"Node readiness status is flapping.\"
\"sum(changes(kube_node_status_condition{status=\"true\\\",condition=\"Ready\\\"}[15m])) by (node) > 2\\n\"
{\"severity\":\"warning\"}
\"KubeletPlegDurationHigh\"
\"Kubelet Pod Lifecycle Event Generator is taking too long to relist.\"
\"node_quantile:kubelet_pleg_relist_duration_seconds:histogram_quantile{quantile=\"0.99\\\"} >= 10\\n\"
{\"severity\":\"warning\"}
\"KubeletPodStartupLatencyHigh\"
\"Kubelet Pod startup latency is too high.\"
\"histogram_quantile(0.99, sum(rate(kubelet_pod_worker_duration_seconds_bucket{job=\"kubelet\\\",
metrics_path=\"/metrics\\\"}[5m])) by (instance, le)) * on(instance) group_left(node)
kubelet_node_name{job=\"kubelet\\\", metrics_path=\"/metrics\\\"} > 60\\n\"
{\"severity\":\"warning\"}
\"KubeletClientCertificateRenewalErrors\"
\"Kubelet has failed to renew its client certificate.\"
\"increase(kubelet_certificate_manager_client_expiration_renew_errors[5m]) > 0\\n\"
{\"severity\":\"warning\"}
\"KubeletServerCertificateRenewalErrors\"
\"Kubelet has failed to renew its server certificate.\"
\"increase(kubelet_server_expiration_renew_errors[5m]) > 0\\n\"
{\"severity\":\"warning\"}
\"KubeletDown\"
\"Target disappeared from Prometheus target discovery.\"
\"absent(up{job=\"kubelet\\\", metrics_path=\"/metrics\\\"} == 1)\\n\"
{\"namespace\":\"kube-system\",\"severity\":\"critical\"}
null
null
\"sum by (cluster, namespace, pod, container) (\\n irate(container_cpu_usage_seconds_total{job=\"kubelet\\\",
metrics_path=\"/metrics/cadvisor\\\", image!=\"\\\"}[5m])\\n) * on (cluster, namespace, pod) group_left(node)
topk by (cluster, namespace, pod) (\\n 1, max by(cluster, namespace, pod, node)
(kube_pod_info{node!=\"\\\"})\\n)\\n\"
null
null
null
\"container_memory_working_set_bytes{job=\"kubelet\\\", metrics_path=\"/metrics/cadvisor\\\", image!=\"\\\"}\\n*
on (namespace, pod) group_left(node) topk by(namespace, pod) (1,\\n max by(namespace, pod, node)
(kube_pod_info{node!=\"\\\"})\\n)\\n\"
null
null
null
\"container_memory_rss{job=\"kubelet\\\", metrics_path=\"/metrics/cadvisor\\\", image!=\"\\\"}\\n* on (namespace,
pod) group_left(node) topk by(namespace, pod) (1,\\n max by(namespace, pod, node)
(kube_pod_info{node!=\"\\\"})\\n)\\n\"
null
null
null
\"container_memory_cache{job=\"kubelet\\\", metrics_path=\"/metrics/cadvisor\\\", image!=\"\\\"}\\n* on
(namespace, pod) group_left(node) topk by(namespace, pod) (1,\\n max by(namespace, pod, node)
(kube_pod_info{node!=\"\\\"})\\n)\\n\"
null
null
null
\"container_memory_swap{job=\"kubelet\\\", metrics_path=\"/metrics/cadvisor\\\", image!=\"\\\"}\\n* on (namespace,
```

```
pod) group_left(node) topk by(namespace, pod) (1,\n  max by(namespace, pod, node)
(kube_pod_info{node!="\n"})\n)\n"
null
null
null
"kube_pod_container_resource_requests{resource="memory",job="kube-state-metrics"} * on (namespace,
pod, cluster)\ngroup_left() max by (namespace, pod, cluster) (\n
(kube_pod_status_phase{phase=~"Pending|Running"} == 1)\n)\n"
null
null
null
"sum by (namespace, cluster) (\n  sum by (namespace, pod, cluster) (\n      max by (namespace, pod,
container, cluster) (\n          kube_pod_container_resource_requests{resource="memory",job="kube-
state-metrics"}\n      ) * on(namespace, pod, cluster) group_left() max by (namespace, pod, cluster)
(\n          kube_pod_status_phase{phase=~"Pending|Running"} == 1\n      )\n  )\n)\n"
null
null
null
"kube_pod_container_resource_requests{resource="cpu",job="kube-state-metrics"} * on (namespace, pod,
cluster)\ngroup_left() max by (namespace, pod, cluster) (\n
(kube_pod_status_phase{phase=~"Pending|Running"} == 1)\n)\n"
null
null
null
"sum by (namespace, cluster) (\n  sum by (namespace, pod, cluster) (\n      max by (namespace, pod,
container, cluster) (\n          kube_pod_container_resource_requests{resource="cpu",job="kube-state-
metrics"}\n      ) * on(namespace, pod, cluster) group_left() max by (namespace, pod, cluster) (\n
kube_pod_status_phase{phase=~"Pending|Running"} == 1\n      )\n  )\n)\n"
null
null
null
"kube_pod_container_resource_limits{resource="memory",job="kube-state-metrics"} * on (namespace, pod,
cluster)\ngroup_left() max by (namespace, pod, cluster) (\n
(kube_pod_status_phase{phase=~"Pending|Running"} == 1)\n)\n"
null
null
null
"sum by (namespace, cluster) (\n  sum by (namespace, pod, cluster) (\n      max by (namespace, pod,
container, cluster) (\n          kube_pod_container_resource_limits{resource="memory",job="kube-state-
metrics"}\n      ) * on(namespace, pod, cluster) group_left() max by (namespace, pod, cluster) (\n
kube_pod_status_phase{phase=~"Pending|Running"} == 1\n      )\n  )\n)\n"
null
null
null
"kube_pod_container_resource_limits{resource="cpu",job="kube-state-metrics"} * on (namespace, pod,
cluster)\ngroup_left() max by (namespace, pod, cluster) (\n
(kube_pod_status_phase{phase=~"Pending|Running"} == 1)\n )\n"
null
null
null
"sum by (namespace, cluster) (\n  sum by (namespace, pod, cluster) (\n      max by (namespace, pod,
container, cluster) (\n          kube_pod_container_resource_limits{resource="cpu",job="kube-state-
metrics"}\n      ) * on(namespace, pod, cluster) group_left() max by (namespace, pod, cluster) (\n
kube_pod_status_phase{phase=~"Pending|Running"} == 1\n      )\n  )\n)\n"
null
null
null
```

```

"max by (cluster, namespace, workload, pod) (\n label_replace(\n label_replace(\n
kube_pod_owner{job=\"kube-state-metrics\", owner_kind=\"ReplicaSet\"},\n      \"replicaset\", \"$1\",
\"owner_name\", \"(.*)\"\\n      ) * on(replicaset, namespace) group_left(owner_name) topk by(replicaset,
namespace) (\n      1, max by (replicaset, namespace, owner_name) (\n
kube_replicaset_owner{job=\"kube-state-metrics\"}\\n      )\\n      ),\n      \"workload\", \"$1\",
\"owner_name\", \"(.*)\"\\n      )\\n)\n"
{"workload_type": "deployment"}
null
null
"max by (cluster, namespace, workload, pod) (\n label_replace(\n kube_pod_owner{job=\"kube-state-
metrics\", owner_kind=\"DaemonSet\"},\n      \"workload\", \"$1\", \"owner_name\", \"(.*)\"\\n      )\\n)\n"
{"workload_type": "daemonset"}
null
null
"max by (cluster, namespace, workload, pod) (\n label_replace(\n kube_pod_owner{job=\"kube-state-
metrics\", owner_kind=\"StatefulSet\"},\n      \"workload\", \"$1\", \"owner_name\", \"(.*)\"\\n      )\\n)\n"
{"workload_type": "statefulset"}
null
null
"histogram_quantile(0.99,
sum(rate(scheduler_e2e_scheduling_duration_seconds_bucket{job=\"scheduler\"}[5m])) without(instance,
pod))\\n"
{"quantile": "0.99"}
null
null
"histogram_quantile(0.99,
sum(rate(scheduler_scheduling_algorithm_duration_seconds_bucket{job=\"scheduler\"}[5m])) without(instance,
pod))\\n"
{"quantile": "0.99"}
null
null
"histogram_quantile(0.99, sum(rate(scheduler_binding_duration_seconds_bucket{job=\"scheduler\"}[5m]))
without(instance, pod))\\n"
{"quantile": "0.99"}
null
null
"histogram_quantile(0.9,
sum(rate(scheduler_e2e_scheduling_duration_seconds_bucket{job=\"scheduler\"}[5m])) without(instance,
pod))\\n"
{"quantile": "0.9"}
null
null
"histogram_quantile(0.9, sum(rate(scheduler_scheduling_algorithm_duration_seconds_bucket{job=\"scheduler\"}[5m]))
without(instance, pod))\\n"
{"quantile": "0.9"}
null
null
"histogram_quantile(0.9, sum(rate(scheduler_binding_duration_seconds_bucket{job=\"scheduler\"}[5m]))
without(instance, pod))\\n"
{"quantile": "0.9"}
null
null
"histogram_quantile(0.5,
sum(rate(scheduler_e2e_scheduling_duration_seconds_bucket{job=\"scheduler\"}[5m])) without(instance,
pod))\\n"
{"quantile": "0.5"}

```

```

null
null
"histogram_quantile(0.5,
sum(rate(scheduler_scheduling_algorithm_duration_seconds_bucket{job=\"scheduler\"}[5m])) without(instance,
pod))\n"
{"quantile": "0.5"}
null
null
"histogram_quantile(0.5, sum(rate(scheduler_binding_duration_seconds_bucket{job=\"scheduler\"}[5m]))
without(instance, pod))\n"
{"quantile": "0.5"}
null
null
"topk by(namespace, pod) (1,\n max by (node, namespace, pod) (\n
label_replace(kube_pod_info{job=\"kube-state-metrics\",node!=\"\"}, \"pod\", \"$1\", \"pod\",
\"(.*)\")\n))\n"
null
null
null
"sum(\n node_memory_MemAvailable_bytes{job=\"node-exporter\"} or\n (\n
node_memory_Buffers_bytes{job=\"node-exporter\"} +\n node_memory_Cached_bytes{job=\"node-exporter\"}
+\n node_memory_MemFree_bytes{job=\"node-exporter\"} +\n node_memory_Slab_bytes{job=\"node-
exporter\"}\n )\n) by (cluster)\n"
null
null
null
"histogram_quantile(0.99, sum(rate(kubelet_pleg_relist_duration_seconds_bucket[5m])) by (instance, le) *
on(instance) group_left(node) kubelet_node_name{job=\"kubelet\", metrics_path=\"/metrics\"})\n"
{"quantile": "0.99"}
null
null
"histogram_quantile(0.9, sum(rate(kubelet_pleg_relist_duration_seconds_bucket[5m])) by (instance, le) *
on(instance) group_left(node) kubelet_node_name{job=\"kubelet\", metrics_path=\"/metrics\"})\n"
{"quantile": "0.9"}
null
null
"histogram_quantile(0.5, sum(rate(kubelet_pleg_relist_duration_seconds_bucket[5m])) by (instance, le) *
on(instance) group_left(node) kubelet_node_name{job=\"kubelet\", metrics_path=\"/metrics\"})\n"
{"quantile": "0.5"}

```

## 10.2.26. openshift-monitoring node-exporter-rules

```

"NodeFilesystemSpaceFillingUp"
"Filesystem is predicted to run out of space within the next 24 hours."
"(\n node_filesystem_avail_bytes{job=\"node-exporter\",fstype!=\"\"} /
node_filesystem_size_bytes{job=\"node-exporter\",fstype!=\"\"} * 100 < 15\nand\n
predict_linear(node_filesystem_avail_bytes{job=\"node-exporter\",fstype!=\"\"}[6h], 24*60*60) < 0\nand\n
node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0)\n"
{"severity": "warning"}
"NodeFilesystemSpaceFillingUp"
"Filesystem is predicted to run out of space within the next 4 hours."
"(\n node_filesystem_avail_bytes{job=\"node-exporter\",fstype!=\"\"} /
node_filesystem_size_bytes{job=\"node-exporter\",fstype!=\"\"} * 100 < 10\nand\n
predict_linear(node_filesystem_avail_bytes{job=\"node-exporter\",fstype!=\"\"}[6h], 4*60*60) < 0\nand\n
node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0)\n"

```



```
{
  "severity": "critical",
  "NodeFilesystemAlmostOutOfSpace": "Filesystem has less than 5% space left.",
  "(\n  node_filesystem_avail_bytes{job=\"node-exporter\",fstype!=\"\"} / node_filesystem_size_bytes{job=\"node-exporter\",fstype!=\"\"} * 100 < 5\nand\n  node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0\n)\n",
  "severity": "warning",
  "NodeFilesystemAlmostOutOfSpace": "Filesystem has less than 3% space left.",
  "(\n  node_filesystem_avail_bytes{job=\"node-exporter\",fstype!=\"\"} / node_filesystem_size_bytes{job=\"node-exporter\",fstype!=\"\"} * 100 < 3\nand\n  node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0\n)\n",
  "severity": "critical",
  "NodeFilesystemFilesFillingUp": "Filesystem is predicted to run out of inodes within the next 24 hours.",
  "(\n  node_filesystem_files_free{job=\"node-exporter\",fstype!=\"\"} / node_filesystem_files{job=\"node-exporter\",fstype!=\"\"} * 100 < 40\nand\n  predict_linear(node_filesystem_files_free{job=\"node-exporter\",fstype!=\"\"}[6h], 24*60*60) < 0\nand\n  node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0\n)\n",
  "severity": "warning",
  "NodeFilesystemFilesFillingUp": "Filesystem is predicted to run out of inodes within the next 4 hours.",
  "(\n  node_filesystem_files_free{job=\"node-exporter\",fstype!=\"\"} / node_filesystem_files{job=\"node-exporter\",fstype!=\"\"} * 100 < 20\nand\n  predict_linear(node_filesystem_files_free{job=\"node-exporter\",fstype!=\"\"}[6h], 4*60*60) < 0\nand\n  node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0\n)\n",
  "severity": "critical",
  "NodeFilesystemAlmostOutOfFiles": "Filesystem has less than 5% inodes left.",
  "(\n  node_filesystem_files_free{job=\"node-exporter\",fstype!=\"\"} / node_filesystem_files{job=\"node-exporter\",fstype!=\"\"} * 100 < 5\nand\n  node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0\n)\n",
  "severity": "warning",
  "NodeFilesystemAlmostOutOfFiles": "Filesystem has less than 3% inodes left.",
  "(\n  node_filesystem_files_free{job=\"node-exporter\",fstype!=\"\"} / node_filesystem_files{job=\"node-exporter\",fstype!=\"\"} * 100 < 3\nand\n  node_filesystem_readonly{job=\"node-exporter\",fstype!=\"\"} == 0\n)\n",
  "severity": "critical",
  "NodeNetworkReceiveErrs": "Network interface is reporting many receive errors.",
  "rate(node_network_receive_errs_total[2m]) / rate(node_network_receive_packets_total[2m]) > 0.01\n",
  "severity": "warning",
  "NodeNetworkTransmitErrs": "Network interface is reporting many transmit errors.",
  "rate(node_network_transmit_errs_total[2m]) / rate(node_network_transmit_packets_total[2m]) > 0.01\n",
  "severity": "warning",
  "NodeHighNumberConntrackEntriesUsed": "Number of conntrack are getting close to the limit.",
  "(node_nf_conntrack_entries / node_nf_conntrack_entries_limit) > 0.75\n",
  "severity": "warning",
  "NodeTextFileCollectorScrapeError": "Node Exporter text file collector failed to scrape.",
  "node_textfile_scrape_error{job=\"node-exporter\"} == 1\n",
  "severity": "warning",
  "NodeClockSkewDetected": "Clock skew detected."
}
```

```
(\n node_timex_offset_seconds > 0.05\nand\n deriv(node_timex_offset_seconds[5m]) >= 0\n)\nor\n(\n
node_timex_offset_seconds < -0.05\nand\n deriv(node_timex_offset_seconds[5m]) <= 0\n)\n"
{"severity":"warning"}
"NodeClockNotSynchronising"
"Clock not synchronising."
"min_over_time(node_timex_sync_status[5m]) == 0\nand\nnode_timex_maxerror_seconds >= 16\n"
{"severity":"warning"}
"NodeRAIDDegraded"
"RAID Array is degraded"
"node_md_disks_required - ignoring (state) (node_md_disks{state=\"active\"}) > 0\n"
{"severity":"critical"}
"NodeRAIDDiskFailure"
"Failed device in RAID array"
"node_md_disks{state=\"failed\"} > 0\n"
{"severity":"warning"}
"NodeFileDescriptorLimit"
"Kernel is predicted to exhaust file descriptors limit soon."
"(\n node_filefd_allocated{job=\"node-exporter\"} * 100 / node_filefd_maximum{job=\"node-exporter\"} >
70\n)\n"
{"severity":"warning"}
"NodeFileDescriptorLimit"
"Kernel is predicted to exhaust file descriptors limit soon."
"(\n node_filefd_allocated{job=\"node-exporter\"} * 100 / node_filefd_maximum{job=\"node-exporter\"} >
90\n)\n"
{"severity":"critical"}
null
null
"count without (cpu, mode) (\n node_cpu_seconds_total{job=\"node-exporter\",mode=\"idle\"})\n"
null
null
null
null
"1 - avg without (cpu) (\n sum without (mode) (rate(node_cpu_seconds_total{job=\"node-exporter\",
mode=~\"idle|iowait|steal\"}[1m]))\n)\n"
null
null
null
"(\n node_load1{job=\"node-exporter\"})\n\n instance:node_num_cpu:sum{job=\"node-exporter\"})\n"
null
null
null
"1 - (\n (\n node_memory_MemAvailable_bytes{job=\"node-exporter\"})\n or\n (\n
node_memory_Buffers_bytes{job=\"node-exporter\"})\n +\n node_memory_Cached_bytes{job=\"node-
exporter\"})\n +\n node_memory_MemFree_bytes{job=\"node-exporter\"})\n +\n
node_memory_Slab_bytes{job=\"node-exporter\"})\n )\n\n node_memory_MemTotal_bytes{job=\"node-
exporter\"})\n)\n"
null
null
null
"rate(node_vmstat_pgmajfault{job=\"node-exporter\"}[1m])\n"
null
null
null
"rate(node_disk_io_time_seconds_total{job=\"node-exporter\",
device=~\"mmcblk.p.+|nvme.+|sd.+|vd.+|xvd.+|dm-.+|dasd.+\"}[1m])\n"
null
null
null
```

```
"rate(node_disk_io_time_weighted_seconds_total{job=\"node-exporter\",
device=~\"mmcblk.p.+|nvme.+|sd.+|vd.+|xvd.+|dm-.+|dasd.+\"}[1m])\n"
null
null
null
"sum without (device) (\n rate(node_network_receive_bytes_total{job=\"node-exporter\",
device!=\"lo\"}[1m])\n)\n"
null
null
null
"sum without (device) (\n rate(node_network_transmit_bytes_total{job=\"node-exporter\",
device!=\"lo\"}[1m])\n)\n"
null
null
null
"sum without (device) (\n rate(node_network_receive_drop_total{job=\"node-exporter\",
device!=\"lo\"}[1m])\n)\n"
null
null
null
"sum without (device) (\n rate(node_network_transmit_drop_total{job=\"node-exporter\",
device!=\"lo\"}[1m])\n)\n"
null
```

## 10.2.27. openshift-monitoring prometheus-k8s-prometheus-rules

```
"PrometheusBadConfig"
"Failed Prometheus configuration reload."
"# Without max_over_time, failed scrapes could create false negatives, see\n#
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for
details.\nmax_over_time(prometheus_config_last_reload_successful{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m]) == 0\n"
{"severity":"warning"}
"PrometheusNotificationQueueRunningFull"
"Prometheus alert notification queue predicted to run full in less than 30m."
"# Without min_over_time, failed scrapes could create false negatives, see\n#
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for details.\n(\n
predict_linear(prometheus_notifications_queue_length{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m], 60 * 30)\n>\n min_over_time(prometheus_notifications_queue_capacity{job=~\"prometheus-
k8s|prometheus-user-workload\"}[5m])\n)\n"
{"severity":"warning"}
"PrometheusErrorSendingAlertsToSomeAlertmanagers"
"Prometheus has encountered more than 1% errors sending alerts to a specific Alertmanager."
"(\n rate(prometheus_notifications_errors_total{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m])\n/\n rate(prometheus_notifications_sent_total{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m])\n)\n* 100\n> 1\n"
{"severity":"warning"}
"PrometheusNotConnectedToAlertmanagers"
"Prometheus is not connected to any Alertmanagers."
"# Without max_over_time, failed scrapes could create false negatives, see\n#
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for
details.\nmax_over_time(prometheus_notifications_alertmanagers_discovered{job=~\"prometheus-
k8s|prometheus-user-workload\"}[5m]) < 1\n"
{"severity":"warning"}
"PrometheusTSDBReloadsFailing"
```

```
"Prometheus has issues reloading blocks from disk."
"increase(prometheus_tsdb_reloads_failures_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[3h]) > 0\n"
{"severity":"warning"}
"PrometheusTSDBCompactionsFailing"
"Prometheus has issues compacting blocks."
"increase(prometheus_tsdb_compactions_failed_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[3h]) > 0\n"
{"severity":"warning"}
"PrometheusNotIngestingSamples"
"Prometheus is not ingesting samples."
"(\n rate(prometheus_tsdb_head_samples_appended_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]) <= 0\nand\n (\n sum without(scrape_job) (prometheus_target_metadata_cache_entries{job=~\"prometheus-k8s|prometheus-user-workload\"}) > 0\n or\n sum without(rule_group) (prometheus_rule_group_rules{job=~\"prometheus-k8s|prometheus-user-workload\"}) > 0\n )\n)\n"
{"severity":"warning"}
"PrometheusDuplicateTimestamps"
"Prometheus is dropping samples with duplicate timestamps."
"rate(prometheus_target_scrapes_sample_duplicate_timestamp_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]) > 0\n"
{"severity":"warning"}
"PrometheusOutOfOrderTimestamps"
"Prometheus drops samples with out-of-order timestamps."
"rate(prometheus_target_scrapes_sample_out_of_order_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]) > 0\n"
{"severity":"warning"}
"PrometheusRemoteStorageFailures"
"Prometheus fails to send samples to remote storage."
"(\n (rate(prometheus_remote_storage_failed_samples_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]) or rate(prometheus_remote_storage_samples_failed_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]))\n/\n (\n (rate(prometheus_remote_storage_failed_samples_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]) or rate(prometheus_remote_storage_samples_failed_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]))\n +\n (rate(prometheus_remote_storage_succeeded_samples_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]) or rate(prometheus_remote_storage_samples_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m]))\n )\n)\n*\n 100\n> 1\n"
{"severity":"warning"}
"PrometheusRemoteWriteBehind"
"Prometheus remote write is behind."
"# Without max_over_time, failed scrapes could create false negatives, see\n#  
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for details.\n(\n max_over_time(prometheus_remote_storage_highest_timestamp_in_seconds{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m])\n- ignoring(remote_name, url) group_right\n max_over_time(prometheus_remote_storage_queue_highest_sent_timestamp_seconds{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m])\n)\n)\n> 120\n"
{"severity":"info"}
"PrometheusRemoteWriteDesiredShards"
"Prometheus remote write desired shards calculation wants to run more than configured max shards."
"# Without max_over_time, failed scrapes could create false negatives, see\n#  
https://www.robustperception.io/alerting-on-gauges-in-prometheus-2-0 for details.\n(\n max_over_time(prometheus_remote_storage_shards_desired{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m])\n)\n>\n max_over_time(prometheus_remote_storage_shards_max{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m])\n)\n)\n"
{"severity":"warning"}
"PrometheusRuleFailures"
```

```
"Prometheus is failing rule evaluations."
"increase(prometheus_rule_evaluation_failures_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[5m])
> 0\n"
{"severity":"warning"}
"PrometheusMissingRuleEvaluations"
"Prometheus is missing rule evaluations due to slow rule group evaluation."
"increase(prometheus_rule_group_iterations_missed_total{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m]) > 0\n"
{"severity":"warning"}
"PrometheusTargetLimitHit"
"Prometheus has dropped targets because some scrape configs have exceeded the targets limit."
"increase(prometheus_target_scrape_pool_exceeded_target_limit_total{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m]) > 0\n"
{"severity":"warning"}
"PrometheusLabelLimitHit"
"Prometheus has dropped targets because some scrape configs have exceeded the labels limit."
"increase(prometheus_target_scrape_pool_exceeded_label_limits_total{job=~\"prometheus-k8s|prometheus-user-
workload\"}[5m]) > 0\n"
{"severity":"warning"}
"PrometheusTargetSyncFailure"
"Prometheus has failed to sync targets."
"increase(prometheus_target_sync_failed_total{job=~\"prometheus-k8s|prometheus-user-workload\"}[30m]) >
0\n"
{"severity":"critical"}
```

## 10.2.28. openshift-monitoring prometheus-k8s-thanos-sidecar-rules

```
"ThanosSidecarBucketOperationsFailed"
"Thanos Sidecar bucket operations are failing"
"sum by (namespace, job, instance)
(rate(thanos_objstore_bucket_operation_failures_total{job=~\"prometheus-(k8s|user-workload)-thanos-
sidecar\"}[5m])) > 0\n"
{"severity":"warning"}
"ThanosSidecarNoConnectionToStartedPrometheus"
"Thanos Sidecar cannot access Prometheus, even though Prometheus seems healthy and has reloaded WAL."
"thanos_sidecar_prometheus_up{job=~\"prometheus-(k8s|user-workload)-thanos-sidecar\"} == 0\nAND on
(namespace, pod)\nprometheus_tsdb_data_replay_duration_seconds != 0\n"
{"severity":"warning"}
```

## 10.2.29. openshift-monitoring prometheus-operator-rules

```
"PrometheusOperatorListErrors"
"Errors while performing list operations in controller."
"(sum by (controller,namespace) (rate(prometheus_operator_list_operations_failed_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[10m])) / sum by (controller,namespace) (rate(prometheus_operator_list_operations_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[10m])))) > 0.4\n"
{"severity":"warning"}
"PrometheusOperatorWatchErrors"
"Errors while performing watch operations in controller."
"(sum by (controller,namespace) (rate(prometheus_operator_watch_operations_failed_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[10m])) / sum by (controller,namespace) (rate(prometheus_operator_watch_operations_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[10m])))) > 0.4\n"
{"severity":"warning"}
"PrometheusOperatorSyncFailed"
"Last controller reconciliation failed"
"min_over_time(prometheus_operator_syncs{status=\"failed\",job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[5m]) > 0\n"
{"severity":"warning"}
"PrometheusOperatorReconcileErrors"
"Errors while reconciling controller."
"(sum by (controller,namespace) (rate(prometheus_operator_reconcile_errors_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[5m])) / (sum by (controller,namespace) (rate(prometheus_operator_reconcile_operations_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[5m])))) > 0.1\n"
{"severity":"warning"}
"PrometheusOperatorNodeLookupErrors"
"Errors while reconciling Prometheus."
"rate(prometheus_operator_node_address_lookup_errors_total{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[5m]) > 0.1\n"
{"severity":"warning"}
"PrometheusOperatorNotReady"
"Prometheus operator not ready"
"min by(namespace, controller) (max_over_time(prometheus_operator_ready{job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[5m]) == 0)\n"
{"severity":"warning"}
"PrometheusOperatorRejectedResources"
"Resources rejected by Prometheus operator"
"min_over_time(prometheus_operator_managed_resources{state=\"rejected\",job=\"prometheus-operator\", namespace=~\"openshift-monitoring|openshift-user-workload-monitoring\"}[5m]) > 0\n"
{"severity":"warning"}
"ConfigReloaderSidecarErrors"
"config-reloader sidecar has not had a successful reload for 10m"
"max_over_time(reloader_last_reload_successful{namespace=~\".+\"}[5m]) == 0\n"
{"severity":"warning"}
```

## 10.2.30. openshift-monitoring telemetry

```

null
null
"count({_name__=~\"cluster:usage:.*|count:up0|count:up1|cluster_version|cluster_version_available_updates
|cluster_operator_up|cluster_operator_conditions|cluster_version_payload|cluster_installer|cluster_infrast
ructure_provider|cluster_feature_set|instance:etcd_object_counts:sum|ALERTS|code:apiserver_request_total:r
ate:sum|cluster:capacity_cpu_cores:sum|cluster:capacity_memory_bytes:sum|cluster:cpu_usage_cores:sum|clust
er:memory_usage_bytes:sum|openshift:cpu_usage_cores:sum|openshift:memory_usage_bytes:sum|workload:cpu_usag
e_cores:sum|workload:memory_usage_bytes:sum|cluster:virt_platform_nodes:sum|cluster:node_instance_type_cou
nt:sum|cnv:vmi_status_running:count|cluster:vmi_request_cpu_cores:sum|node_role_os_version_machine:cpu_cap
acity_cores:sum|node_role_os_version_machine:cpu_capacity_sockets:sum|subscription_sync_total|olm_resoluti
on_duration_seconds|csv_succeeded|csv_abnormal|cluster:kube_persistentvolumeclaim_resource_requests_storag
e_bytes:provisioner:sum|cluster:kubelet_volume_stats_used_bytes:provisioner:sum|ceph_cluster_total_bytes|c
eph_cluster_total_used_raw_bytes|ceph_health_status|job:ceph_osd_metadata:count|job:kube_pv:count|job:ceph
_pools_iops:total|job:ceph_pools_iops_bytes:total|job:ceph_versions_running:count|job:noobaa_total_unhealt
hy_buckets:sum|job:noobaa_bucket_count:sum|job:noobaa_total_object_count:sum|noobaa_accounts_num|noobaa_to
tal_usage|console_url|cluster:network_attachment_definition_instances:max|cluster:network_attachment_defin
ition_enabled_instance_up:max|cluster:ingress_controller_aws_nlb_active:sum|insightsclient_request_send_to
tal|cam_app_workload_migrations|cluster:apiserver_current_inflight_requests:sum|max_over_time:2m|cluster:a
lertmanager_integrations:max|cluster:telemetry_selected_series:count|openshift:prometheus_tsdb_head_series
:sum|openshift:prometheus_tsdb_head_samples_appended_total:sum|monitoring:container_memory_working_set_byt
es:sum|namespace_job:scrape_series_added:topk3_sum1h|namespace_job:scrape_samples_post_metric_relabeling:t
opk3|monitoring:haproxy_server_http_responses_total:sum|rhmi_status|cluster_legacy_scheduler_policy|cluste
r_master_schedulable|che_workspace_status|che_workspace_started_total|che_workspace_failure_total|che_work
space_start_time_seconds_sum|che_workspace_start_time_seconds_count|cco_credentials_mode|cluster:kube_pers
istentvolume_plugin_type_counts:sum|visual_web_terminal_sessions_total|acm_managed_cluster_info|cluster:vs
phere_vcenter_info:sum|cluster:vsphere_esxi_version_total:sum|cluster:vsphere_node_hw_version_total:sum|op
enshift:build_by_strategy:sum|rhods_aggregate_availability|rhods_total_users|instance:etcd_disk_wal_fsync_
duration_seconds:histogram_quantile|instance:etcd_mvcc_db_total_size_in_bytes:sum|instance:etcd_network_pe
er_round_trip_time_seconds:histogram_quantile|instance:etcd_mvcc_db_total_size_in_use_in_bytes:sum|instanc
e:etcd_disk_backend_commit_duration_seconds:histogram_quantile|jaeger_operator_instances_storage_types|jae
ger_operator_instances_strategies|jaeger_operator_instances_agent_strategies|apps_vcs:cores_by_product:sum|
nto_custom_profiles:count|openshift_csi_share_configmap|openshift_csi_share_secret|openshift_csi_share_mou
nt_failures_total|openshift_csi_share_mount_requests_total\",alertstate=~\"firing|\",quantile=~\"0.99|0.99
|0.99|\"}})"
null

```

### 10.2.31. openshift-monitoring thanos-querier

```
"ThanosQueryHttpRequestQueryErrorRateHigh"
"Thanos Query is failing to handle requests."
"(\n sum by (namespace, job) (rate(http_requests_total{code=~\"5..\", job=\"thanos-querier\",
handler=\"query\"}[5m]))\n/\n sum by (namespace, job) (rate(http_requests_total{job=\"thanos-querier\",
handler=\"query\"}[5m]))\n) * 100 > 5\n"
{"severity":"warning"}
"ThanosQueryHttpRequestQueryRangeErrorRateHigh"
"Thanos Query is failing to handle requests."
"(\n sum by (namespace, job) (rate(http_requests_total{code=~\"5..\", job=\"thanos-querier\",
handler=\"query_range\"}[5m]))\n/\n sum by (namespace, job) (rate(http_requests_total{job=\"thanos-querier\",
handler=\"query_range\"}[5m]))\n) * 100 > 5\n"
{"severity":"warning"}
"ThanosQueryGrpcServerErrorRate"
"Thanos Query is failing to handle requests."
"(\n sum by (namespace, job)
(rate(grpc_server_handled_total{grpc_code=~\"Unknown|ResourceExhausted|Internal|Unavailable|DataLoss|DeadlineExceeded\", job=\"thanos-querier\"}[5m]))\n/\n sum by (namespace, job)
(rate(grpc_server_started_total{job=\"thanos-querier\"}[5m]))\n) * 100 > 5\n)\n"
{"severity":"warning"}
"ThanosQueryGrpcClientErrorRate"
"Thanos Query is failing to send requests."
"(\n sum by (namespace, job) (rate(grpc_client_handled_total{grpc_code!=\"OK\", job=\"thanos-querier\"}[5m]))\n/\n sum by (namespace, job) (rate(grpc_client_started_total{job=\"thanos-querier\"}[5m]))\n) * 100 > 5\n"
{"severity":"warning"}
"ThanosQueryHighDNSFailures"
"Thanos Query is having high number of DNS failures."
"(\n sum by (namespace, job) (rate(thanos_query_store_apis_dns_failures_total{job=\"thanos-querier\"}[5m]))\n/\n sum by (namespace, job)
(rate(thanos_query_store_apis_dns_lookups_total{job=\"thanos-querier\"}[5m]))\n) * 100 > 1\n"
{"severity":"warning"}
```

### 10.2.32. openshift-multus prometheus-k8s-rules

```
null
null
"max (network_attachment_definition_enabled_instance_up) by (networks)\n"
null
null
null
"max (network_attachment_definition_instances) by (networks)\n"
null
```

### 10.2.33. openshift-operator-lifecycle-manager olm-alert-rules



```
"CsvAbnormalFailedOver2Min"
"CSV failed for over 2 minutes"
"csv_abnormal{phase=~\"^Failed$\"}"
{"namespace":{"{{ $labels.namespace }}"}, "severity": "warning"}
"CsvAbnormalOver30Min"
"CSV abnormal for over 30 minutes"
"csv_abnormal{phase=~\"(^Replacing$|^Pending$|^Deleting$|^Unknown$)\"}"
{"namespace":{"{{ $labels.namespace }}"}, "severity": "warning"}
"InstallPlanStepAppliedWithWarnings"
"API returned a warning when modifying an operator"
"sum(increase(installplan_warnings_total[5m])) > 0"
{"severity": "warning"}
```

### 10.2.34. openshift-sdn networking-rules

```
"NodeWithoutSDNController"
"All control plane nodes should be running a sdn controller pod, {{ $labels.node }} is not."
"count(kube_node_role{role=\"master\"} == 1) != count(kube_pod_info{namespace=\"openshift-sdn\",
pod=~\"sdn-controller.*\"})\n"
{"severity": "warning"}
"NodeWithoutSDNPod"
"All nodes should be running an sdn pod, {{ $labels.node }} is not."
"(kube_node_info unless on(node) topk by (node) (1, kube_pod_info{namespace=\"openshift-sdn\",
pod=~\"sdn-[^-]*\"})) > 0\n"
{"severity": "warning"}
"NodeProxyApplySlow"
"SDN pod {{ $labels.pod }} on node {{ $labels.node }} is taking too long, on average, to apply kubernetes
service rules to iptables."
"histogram_quantile(.95, kubeproxy_sync_proxy_rules_duration_seconds_bucket)\n* on(namespace, pod)
group_right topk by (namespace, pod) (1, kube_pod_info{namespace=\"openshift-sdn\", pod=~\"sdn-[^-]*\"})
> 15\n"
{"severity": "warning"}
"ClusterProxyApplySlow"
"The cluster is taking too long, on average, to apply kubernetes service rules to iptables."
"histogram_quantile(0.95, sum(rate(kubeproxy_sync_proxy_rules_duration_seconds_bucket[5m])) by (le)) >
10\n"
{"severity": "warning"}
"NodeProxyApplyStale"
"SDN pod {{ $labels.pod }} on node {{ $labels.node }} has stale kubernetes service rules in iptables."
"(kubeproxy_sync_proxy_rules_last_queued_timestamp_seconds -
kubeproxy_sync_proxy_rules_last_timestamp_seconds)\n* on(namespace, pod) group_right() topk by (namespace,
pod) (1, kube_pod_info{namespace=\"openshift-sdn\", pod=~\"sdn-[^-]*\"})\n> 30\n"
{"severity": "warning"}
"SDNPodNotReady"
"SDN pod {{ $labels.pod }} on node {{ $labels.node }} is not ready."
"kube_pod_status_ready{namespace='openshift-sdn', condition='true'} == 0\n"
{"severity": "warning"}
```

## 10.3. Prometheus Context

```
{ "fsGroup": 65534, "runAsNonRoot": true, "runAsUser": 65534 }
"15d"
{ "requests": { "cpu": "70m", "memory": "1Gi" } }
```

## 10.4. Service Monitors

NAMESPACE	NAME
openshift-apiserver-operator	openshift-apiserver-operator
openshift-apiserver	openshift-apiserver
openshift-apiserver	openshift-apiserver-operator-check-endpoints
openshift-authentication-operator	authentication-operator
openshift-authentication	oauth-openshift
openshift-cloud-credential-operator	cloud-credential-operator
openshift-cluster-machine-approver	cluster-machine-approver
openshift-cluster-node-tuning-operator	node-tuning-operator
openshift-cluster-samples-operator	cluster-samples-operator
openshift-cluster-storage-operator	cluster-storage-operator
openshift-cluster-version	cluster-version-operator
openshift-config-operator	config-operator
openshift-console-operator	console-operator
openshift-console	console
openshift-controller-manager-operator	openshift-controller-manager-operator
openshift-controller-manager	openshift-controller-manager
openshift-dns-operator	dns-operator
openshift-dns	dns-default
openshift-etcd-operator	etcd-operator
openshift-image-registry	image-registry
openshift-image-registry	image-registry-operator
openshift-ingress-operator	ingress-operator
openshift-ingress	router-default
openshift-insights	insights-operator
openshift-kube-apiserver-operator	kube-apiserver-operator
openshift-kube-apiserver	kube-apiserver
openshift-kube-controller-manager-operator	kube-controller-manager-operator
openshift-kube-controller-manager	kube-controller-manager
openshift-kube-scheduler-operator	kube-scheduler-operator
openshift-kube-scheduler	kube-scheduler
openshift-machine-api	cluster-autoscaler-operator
openshift-machine-api	machine-api-controllers
openshift-machine-api	machine-api-operator
openshift-machine-config-operator	machine-config-daemon
openshift-marketplace	marketplace-operator
openshift-monitoring	alertmanager-main
openshift-monitoring	cluster-monitoring-operator
openshift-monitoring	etcd
openshift-monitoring	grafana
openshift-monitoring	kube-state-metrics
openshift-monitoring	kubelet
openshift-monitoring	node-exporter
openshift-monitoring	openshift-state-metrics
openshift-monitoring	prometheus-adapter
openshift-monitoring	prometheus-k8s
openshift-monitoring	prometheus-operator

openshift-monitoring	telemeter-client
openshift-monitoring	thanos-querier
openshift-monitoring	thanos-sidecar
openshift-multus	monitor-multus-admission-controller
openshift-multus	monitor-network
openshift-network-diagnostics	network-check-source
openshift-oauth-apiserver	openshift-oauth-apiserver
openshift-operator-lifecycle-manager	catalog-operator
openshift-operator-lifecycle-manager	olm-operator
openshift-sdn	monitor-sdn
openshift-sdn	monitor-sdn-controller
openshift-service-ca-operator	service-ca-operator

## 10.5. Pod Monitors

## 10.6. Alert Managers

NAMESPACE	NAME	VERSION	REPLICAS	AGE
openshift-monitoring	main	0.23.0	1	19h

## 10.7. Monitoring Agents & Dashboards

Agent	Status
cluster-monitoring-operator	OK
kube-state-metrics	OK
openshift-state-metrics	OK
node-exporter	OK
thanos-querier	OK
grafana	OK
telemeter-client	OK

# Chapter 11. operators

## 11.1. Degraded Cluster Operators

## 11.2. Unavailable Cluster Operators

## 11.3. Cluster Services Versions

NAMESPACE	NAME	DISPLAY	VERSION	REPLACES	PHASE
openshift-operator-lifecycle-manager	packageserver	Package Server	0.19.0		Succeeded

## 11.4. Operator Groups

NAMESPACE	NAME	AGE
openshift-monitoring	openshift-cluster-monitoring	41d
openshift-operator-lifecycle-manager	olm-operators	41d
openshift-operators	global-operators	41d

## 11.5. Operator Sources

## **Chapter 12. Service Mesh**

### **12.1. Service Mesh ControlPlane**

### **12.2. Service Mesh Members**

### **12.3. Service Mesh Member Rolls**

## Chapter 13. Applications

### 13.1. Non-Ready Deployments

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
ocapp-test	my-app2	0/1	0	0	13d
ocapp-test	my-app3	0/1	0	0	13d

### 13.2. Unavailable Deployments

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
ocapp-test	my-app2	0/1	0	0	13d
ocapp-test	my-app3	0/1	0	0	13d

### 13.3. Failed Builds

### 13.4. Inactive projects